



Scalable Internet Architectures

how to build scalable production Internet services and...

how not to build them

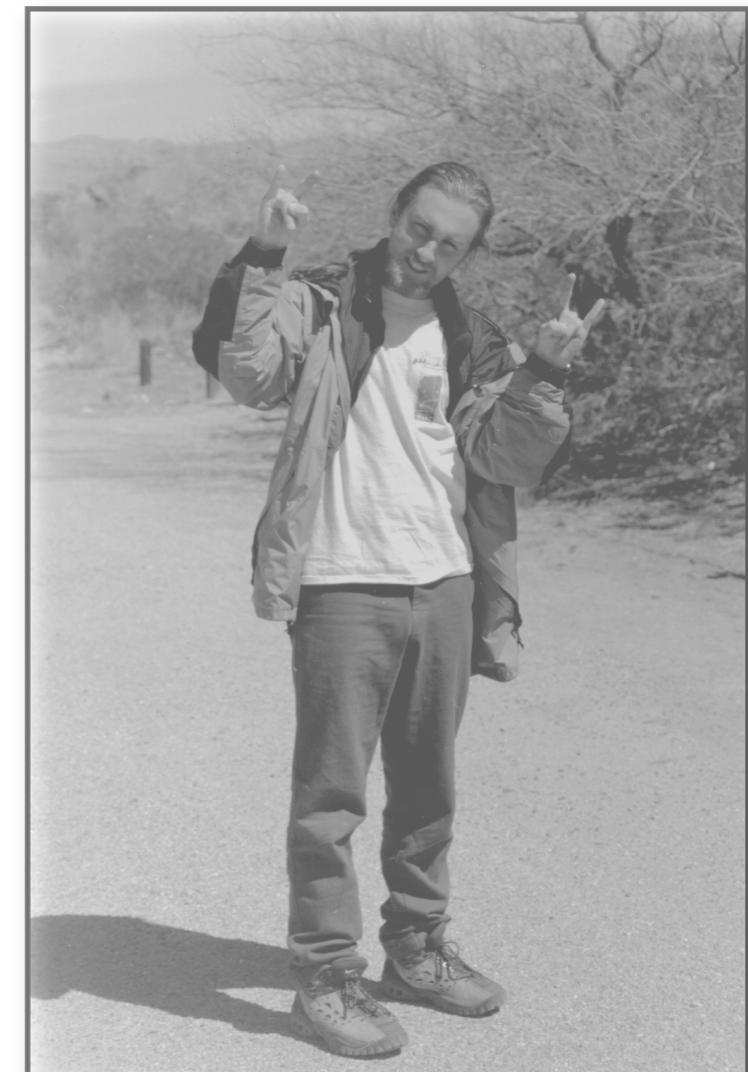




A bit about the speaker

Principal Consultant
OmniTI Computer Consulting, Inc.

- **open-source developer**
 - mod_backhand, wackamole, Daiquiri, Spread, OpenSSH/SecurID, a variety of CPAN modules, etc.
- **closed-source developer**
 - Ecelerity (MTA), EC Cluster (MTA Clustering)
- **open-source advocate**
 - Closed source software has technical risk.
- **closed-source advocate**
 - It's about business, *not* software. Finding the right tool for the job sometimes leads to closed-source solutions.



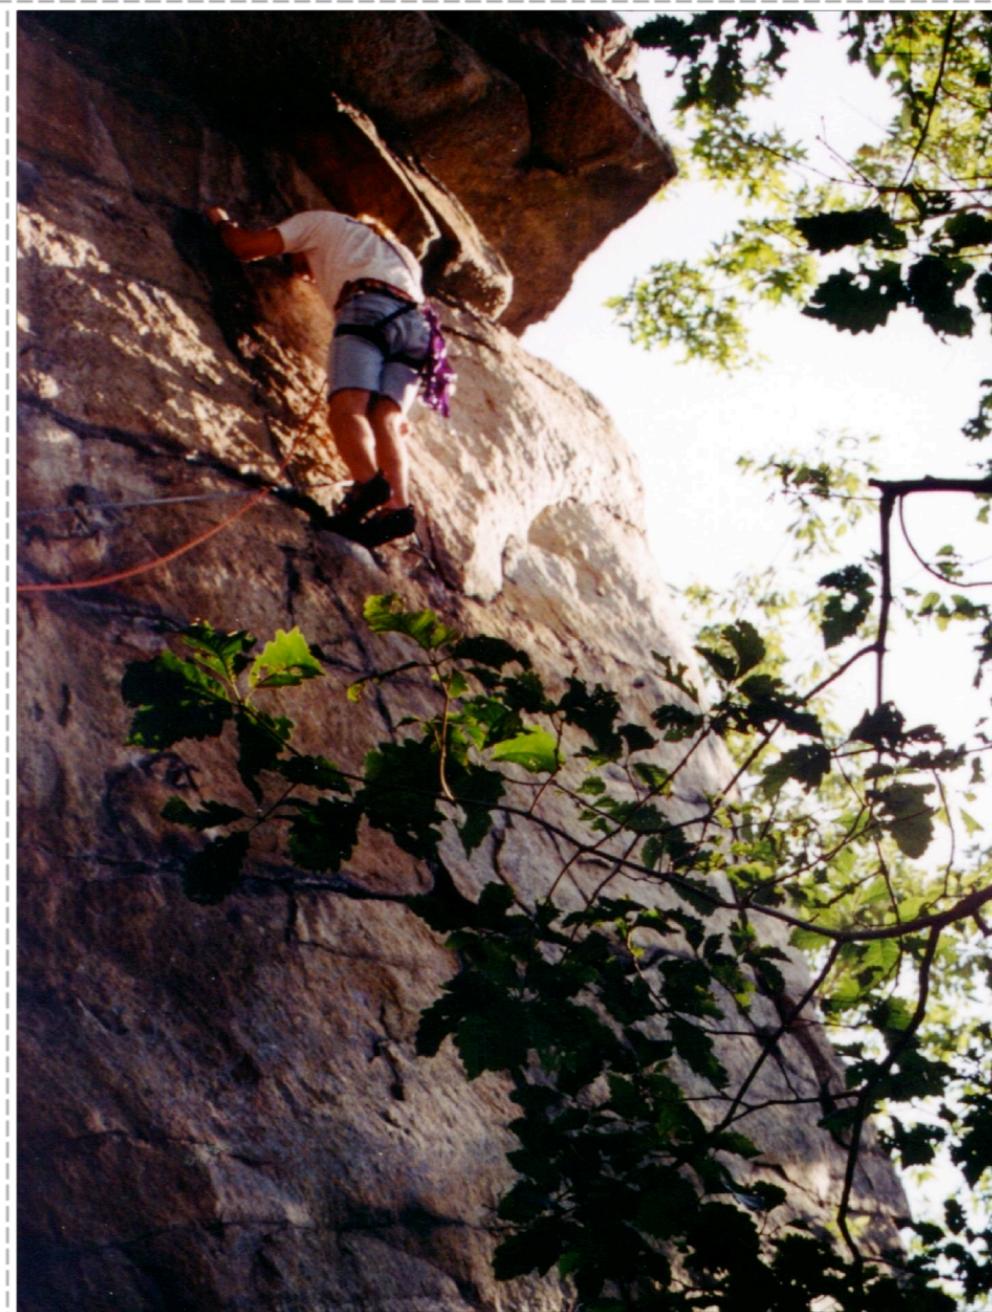
What is Scalability?

Definition:

How well a solution to some problem will work when the size of the problem increases.

What's missing?

**... when the size decreases
the solution must fit**





Production Environments

high uptime

low maintenance

formal procedures

cost controlled



High Uptime

Availability despite individual system failures



parallel servers

- all servers are live and can handle transactions
 - cheap and common for web servers
 - expensive for databases



hot spare/standby

- fail-over system that is seamless and immediate (automated)
 - common for HA/LB solutions
 - many databases have built-in facilities providing hot-spare service



warm spare/standby

- fail-over system is nearly immediate, but not seamless (not automated)
 - common technique for databases, cheap and easy



cold spare/standby

- "I have the equipment and backups to get it running if it were to fail."



Maintenance

The single largest expense in most environments

● Contributing factors:

- The number of unique required products in the architecture
- The stability and “replaceability” of required products
- Uneducated development and implementation decisions
- The complexity and frequency of staging and pushing new code





Formal Procedures

**“Scalability marginally impacts procedure
Procedure grossly impacts scalability”**

- developer code review
- religious use of revision control
- planned and reviewed upgrade strategies
- intelligent, low-cost (resources) push procedures





Three Simple Rules

optimize where it counts

complexity has costs

use the right tool





Three Simple Rules

#1: Amdahl's Law

Good

improving execution time by 50%
of code that executes 2% of the time
results in 1% performance improvement

Better

improving execution time by 10%
of code that executes 80% of the time
results in 8% performance improvement





Three Simple Rules

#2: Complex architectures are expensive



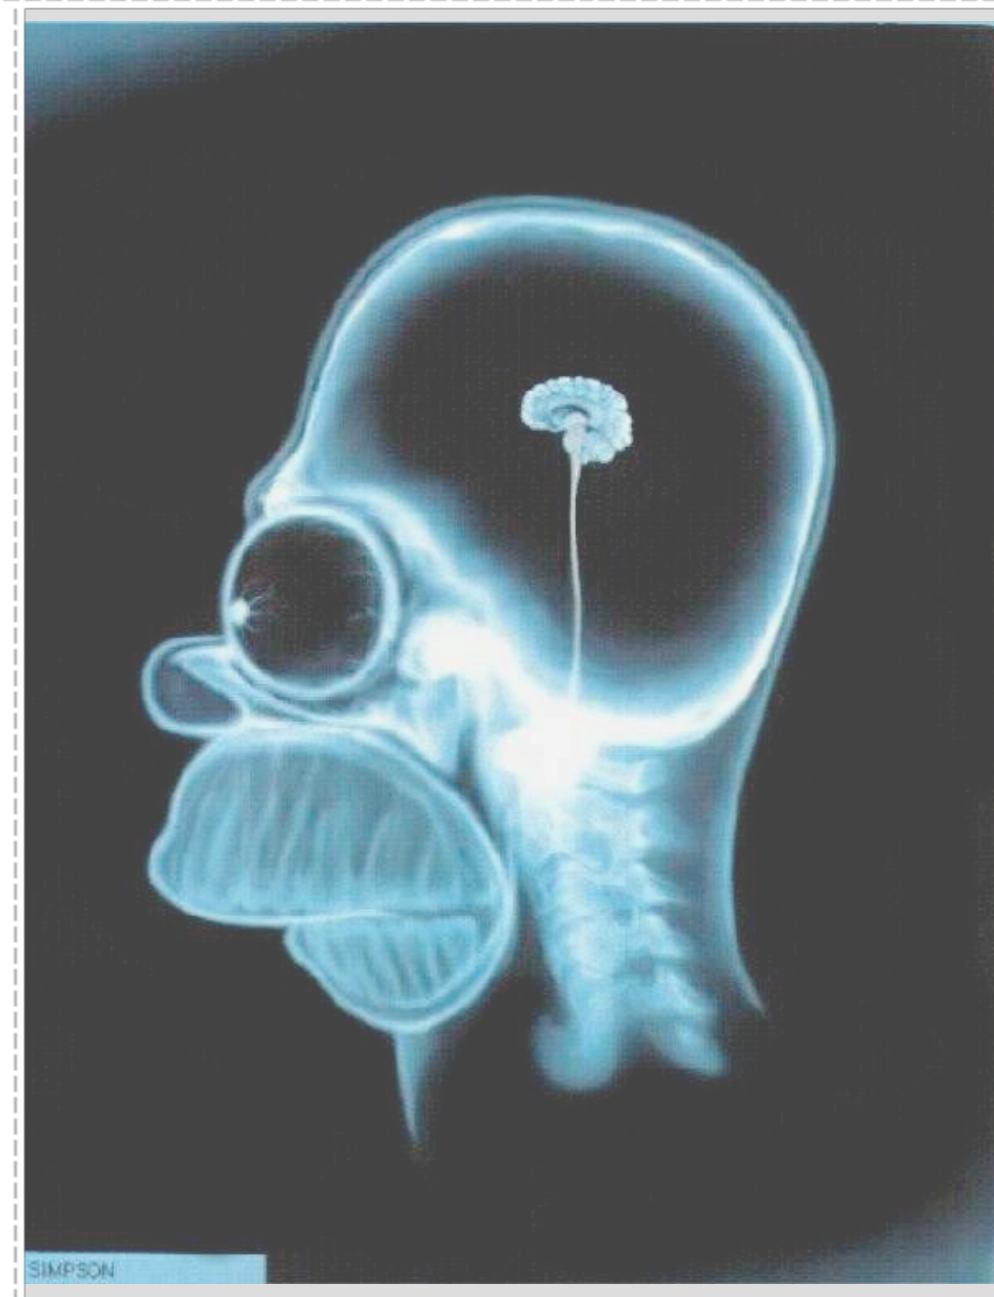
- adding an additional architectural component to a service or set of services increases the system complexity linearly
- requiring an additional architectural component for a service increases the system complexity exponentially



Three Simple Rules

#3: Using the wrong tool is expensive (and stupid)

- using a tool because it is easy or familiar doesn't make it right
- it is often a gratuitous waste of resources
- white papers are marketing tools and may not represent the most practical solution
- it's about good design and implementation practices





Building Production Systems



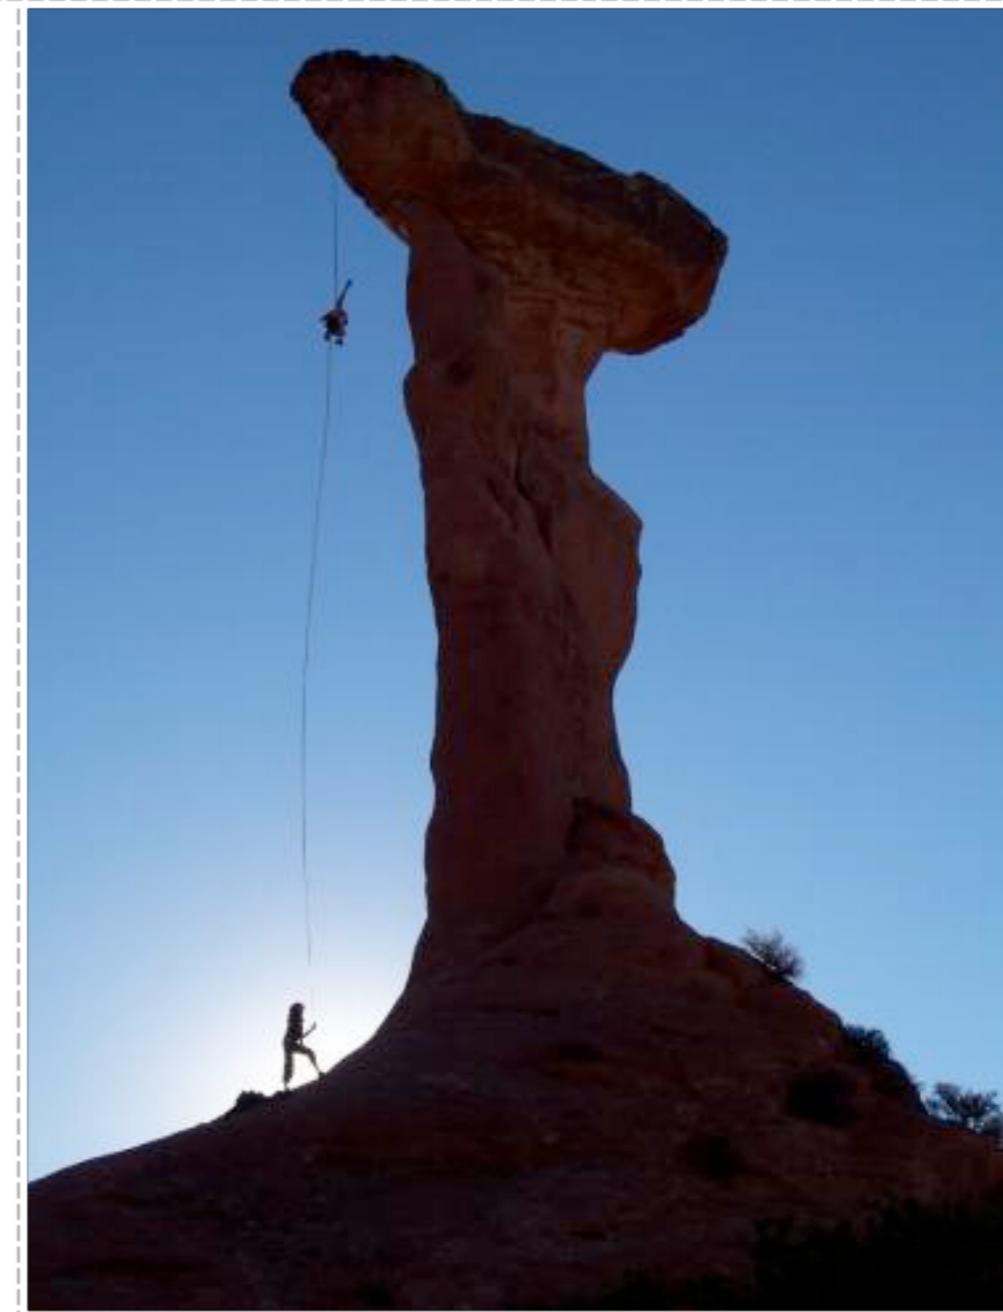
Production Fundamentals

- understand the stability of the software
- understand the velocity of development
- understand administrative aspects
- understand the likelihood of failure
and the support for each component



Software Stability

- **Stability is not just reliability**
- **Also consider:**
 - release cycles
 - upgrade paths
 - feature additions,
“deprecations”, and removals



The Need For Speed

vs.

The Need For Control

- Understand the velocity of development
 - For Small Projects: use revision control
 - For Large Projects: use revision control
 - For ALL Projects: use revision control
- No revision control?
- Accident waiting to happen



The Need For Speed

vs.

The Need For Control

- Unchecked speed is costly
- Rapid release cycles (once/day) are needed in some businesses
- An equilibrium must be achieved or the situation will explode
- Properly used revision control allows for speed and control
- It is challenging, but meticulous unwavering adherence to policy and procedure will deliver you from disaster.





Administration

● This deserves a lot of attention
(despite the single slide here)

● Systems Administration costs money

- Short release cycles on components means perpetual administration
- Constant change in development product results in different stress on:
 - Databases, Networks, Systems... and the people that maintain them
- Adding components or complicating the architecture complicates:
 - Monitoring
 - Upgrading
 - Scaling down should the need arise



Likelihood of Failure (the hidden administrative nightmare)

- ➊ Internally Developed Application Failures Suck.
- ➋ Third-Party Component Failures Suck More!
 - ➌ It is seen as an administration responsibility
 - ➌ Regardless if developers dictated their inclusion in the architecture
 - ➌ SAs, NAs, and DBAs suddenly become responsible for the ongoing maintenance of all third-party products -- open source or commercial
 - ➌ This is often beyond the expertise/attention of the individual or team
 - ➌ Systems fail, it's part of life
 - ➌ Chronic problems and failures will explode your TCO



Likelihood of Failure (solution to the hidden administrative nightmare)

- **Don't leave “requirement” assessments at:**
 - “This won’t work... but you’re the boss”
- **Worst**
 - implementing something that won’t work
 - being responsible for making it work
 - getting fired for perceived incompetence
- **Bad**
 - getting fired for refusing to implement something that has no hope of working
- **Best**
 - work with the development team to revise requirements and architectural needs





Clustered Image Serving





Goal

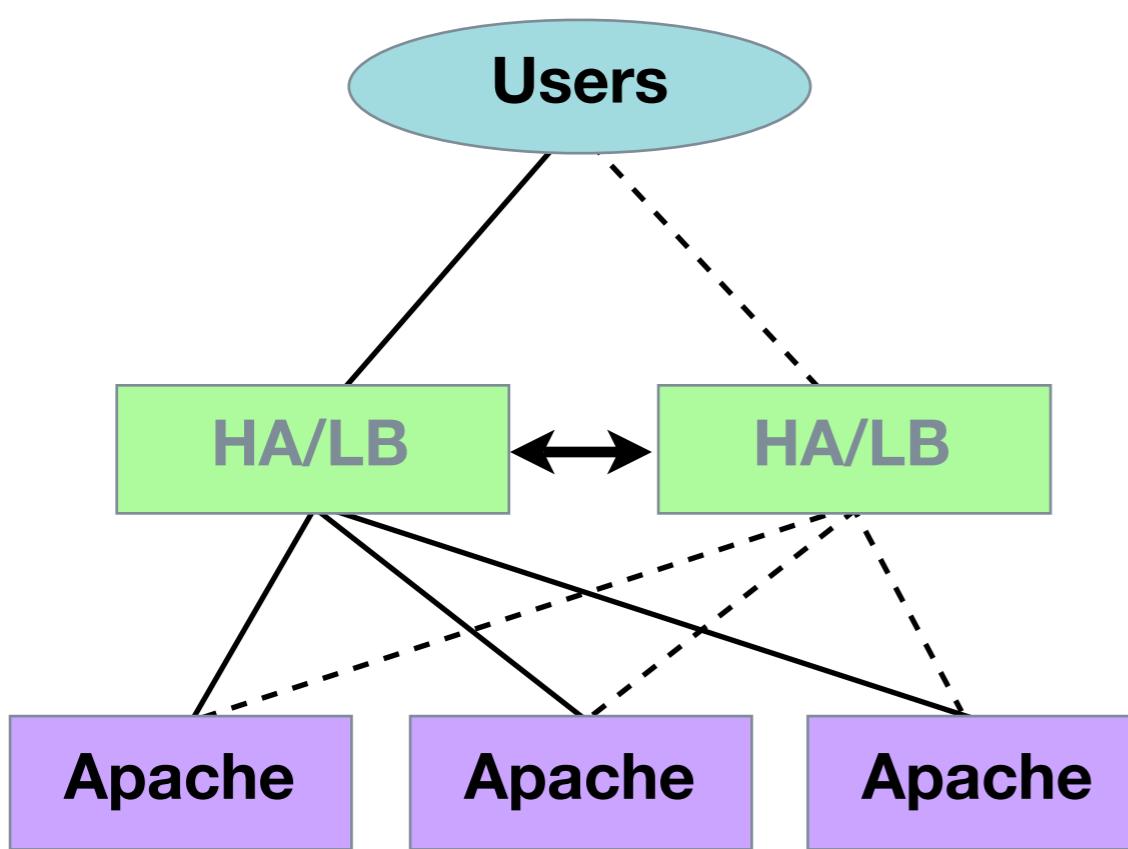
- Static image serving
- 120MBs throughput
- 24x7 uptime requirements
- Three geographically distributed sites





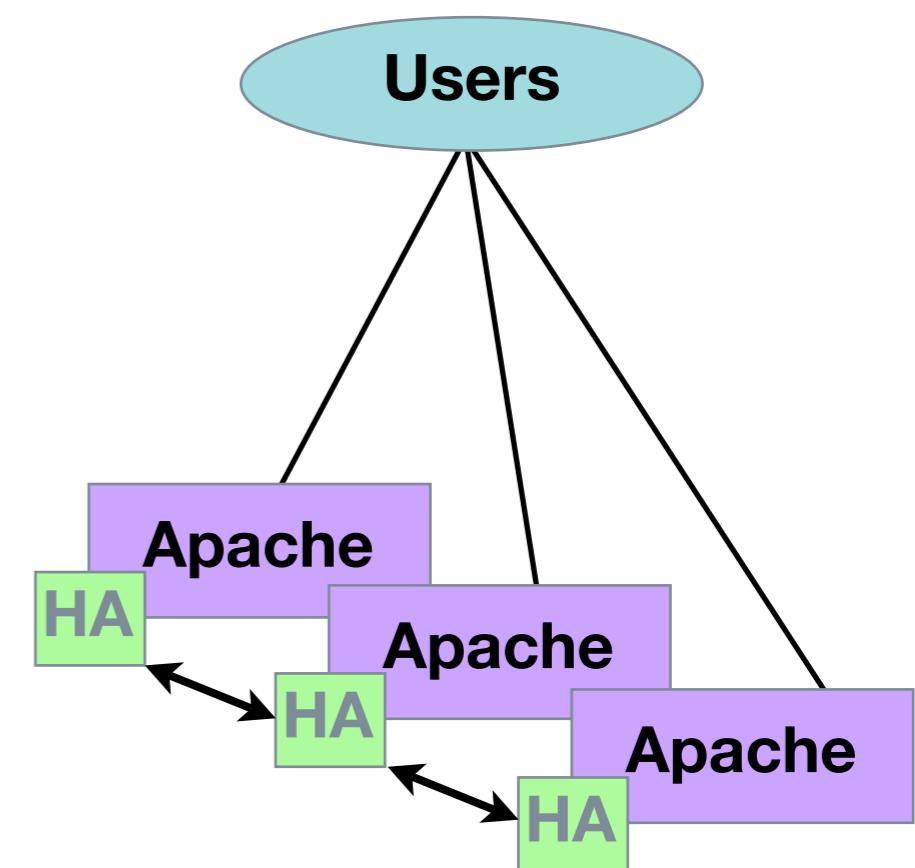
Decisions

high uptime



“White Paper” Approach

expensive, dedicated, single-purpose
HA/LB devices



Peer-based HA

cheap and reusable commodity machines



The ~~Tired~~ Tiered Approach

● Pros:

- Fine-grained, connection-based request distribution (load balancing)
- 100,000+ concurrent connections
- Session management (sticky)
- One IP per service

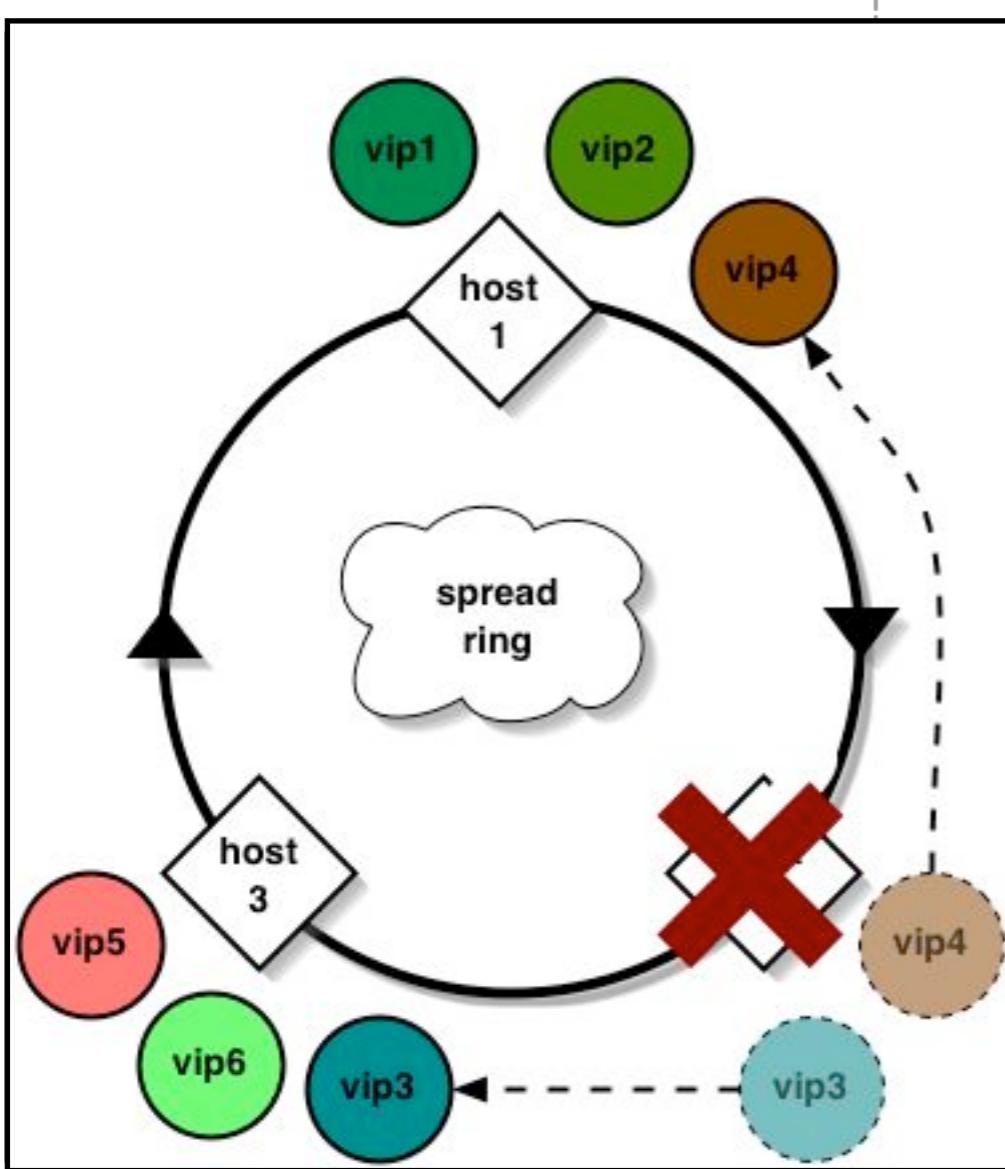
● Cons:

- Expensive
- Single purpose
- Your HA solution needs HA!!!
- 3 locations requires 6 units
- High maintenance
(additional hardware component)



Peer-based HA

Wackamole



Pros:

- No specialized hardware
- Low maintenance (software daemon)
- Simple
- Free

Cons:

- Naïve load balancing (DNS RR)
- Requires multiple IPs for a single service (bad for multi-SSL)



Policy & Procedure

- Pushing content
 - Even for small (~100Mb) image repositories, pushes are expensive
 - dumb protocols have horrible network costs
 - rsync still incurs substantial I/O for each “mirror”
 - multicast rsync could work, but there are no solid implementations

- Pulling content
 - Assuming a slow rate of change, cache-on-demand is solid
 - Use Apache + mod_proxy (Reverse Proxy + Caching)
 - Fine-grained cache purging is a challenge



Scaling Up

3 Sites

- Goal
 - 200Mbs throughput requirement
 - The goal is lower latency
 - only 2 web servers per site needed for fault tolerance
- Traditional “White Paper” Approach
 - 3 x dual HA/LB
 - 3 x 2 image web servers
$$\begin{array}{r} 3 \times 2 \times \$10000 \\ + 3 \times 2 \times \$2000 \\ \hline \$72000 \end{array}$$
- Peer-based HA Solution
 - 3 x 2 image web servers
$$\begin{array}{r} 3 \times 2 \times \$2000 \\ \hline \$12000 \end{array}$$



Scaling Down

1 Site

- Goal

- 10Mbs throughput requirement
- The goal is lower latency
- only 2 web servers per site needed for fault tolerance

- Traditional “White Paper” Approach

- dual HA/LB
- 2 image web servers

2 x \$10000

+ 2 x \$2000

\$24000

- Peer-based HA Solution

- 2 image web servers

2 x \$2000

\$4000



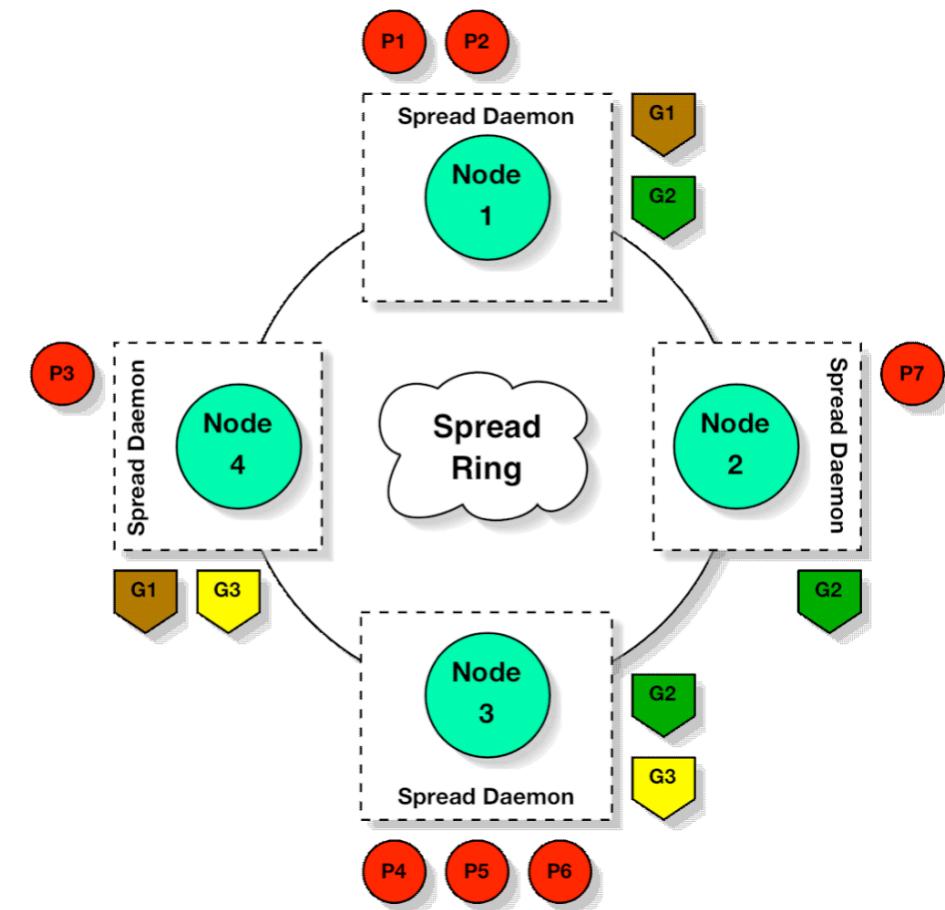
Technical Details

- **Each box running FreeBSD 5-stable**
 - <http://www.freebsd.org/>
- **Spread v3.17.3**
 - <http://www.spread.org/>
- **wackamole 2.1.2**
 - <http://www.backhand.org/wackamole/>
- **Apache 1.3.33/mod_ssl + mod_proxy + patches**
 - <http://www.apache.org/>
 - <http://www.omniti.com/~george/>



Spread: What is it?

- Group Communication
 - Messaging Bus
 - Membership
- Clear Delivery Semantics
 - Reliable or Unreliable
 - FIFO, Causal
 - Agreed, Safe
 - View of membership for delivery
- Fast and Efficient
 - N subscribers $\neq N \times$ bandwidth
 - Multicast or broadcast
- Usable
 - C, Perl, Python, Java, PHP, Ruby API



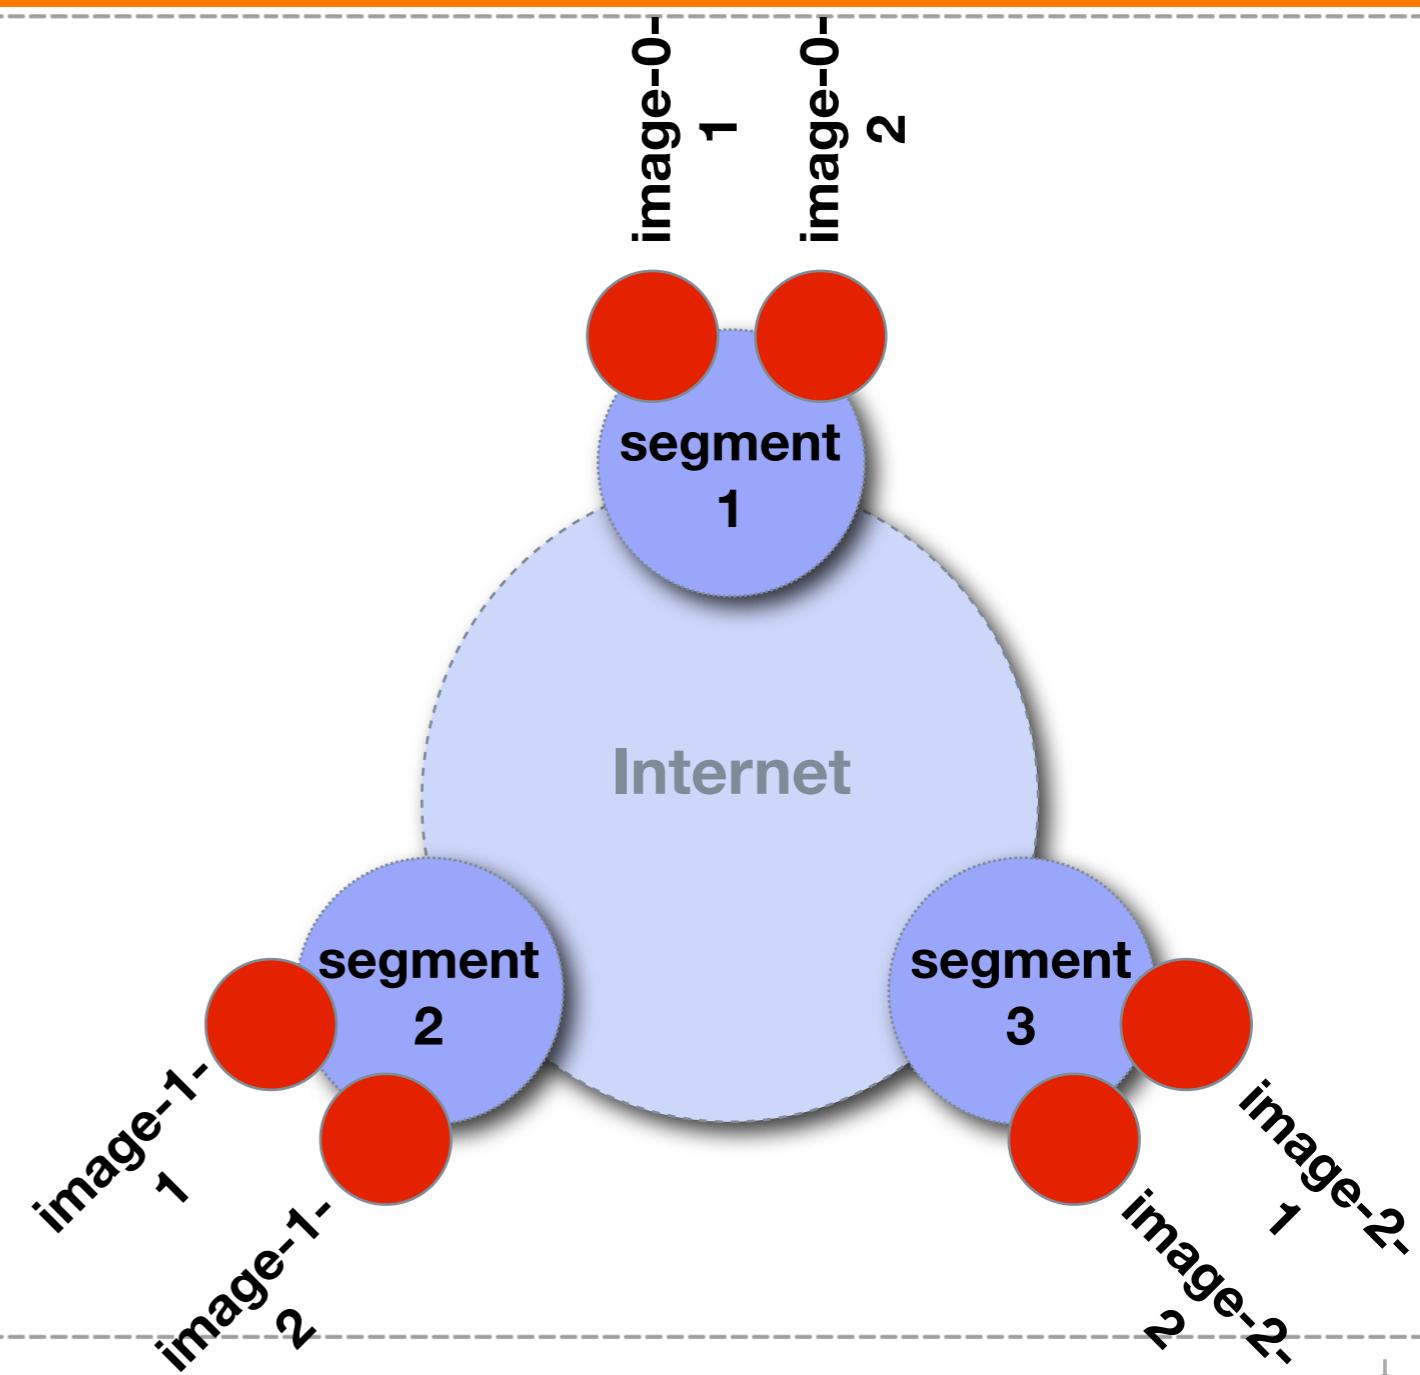


Spread: Configuration

```
# New Jersey Site
Spread_Segment 225.0.1.1:4803 {
    image-0-1          a.b.c.101
    image-0-2          a.b.c.102
}

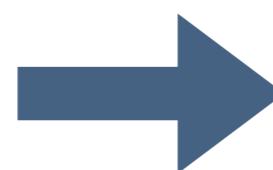
# San Jose Site
Spread_Segment 225.0.1.2:4803 {
    image-1-1          d.e.f.101
    image-1-2          d.e.f.102
}

# Germany Site
Spread_Segment 225.0.1.3:4803 {
    image-2-1          g.h.i.101
    image-2-2          g.h.i.102
}
```



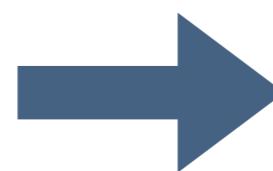
Wackamole: Configuration

Spread Daemon & Group



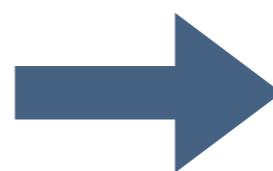
```
Spread = 4803  
SpreadRetryInterval = 5s  
Group = wack1  
Control = /var/run/wack.it
```

Virtual Interfaces Controlled



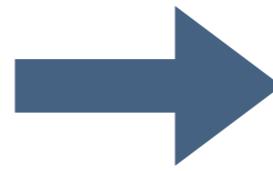
```
Prefer None  
VirtualInterfaces {  
    { fxp0:a.b.c.111/32 }  
    { fxp0:a.b.c.112/32 }  
}
```

Notifications of ownership



```
Arp-Cache = 90s  
Notify {  
    fxp0:a.b.c.1/32  
    fxp0:a.b.c.0/24 throttle 16  
    arp-cache  
}
```

Balancing parameters



```
balance {  
    AcquisitionsPerRound = all  
    interval = 4s  
}  
mature = 5s
```

Apache: Configuration

```
# Don't act as a free image caching service!
```

```
<Directory proxy:>
    deny from all
</Directory>
```

```
# But act provide service to us
```

```
<Directory proxy:http://www.example.com/*>
    allow from all
</Directory>
```

```
RewriteEngine on
```

```
RewriteLogLevel 0
```

```
RewriteRule ^proxy: - [F]
```

```
RewriteRule ^http|ftp: - [F]
```

```
RewriteRule ^V*([^\V]+)(.*)$ http://$1$2 [P,L]
```

```
RewriteRule .* - [F]
```

```
ProxyRequests on
```

```
CacheRoot /data/cache
```

```
CacheSize 5120000
```

```
ProxyPassReverse / http://www.example.com/
```



Why patch mod_proxy?

- **mod_proxy hashes URLs for local caching**
 - better distribution of files over directories
 - nice to your filesystem
 - good for forward caches
 - makes purging individual URLs less intuitive

- **Patched to write intuitive filenames**
 - a URL like: `http://www.example.com/logo.gif` becomes `/data/cache/www.example.com/logo.gif`
 - SAs can troubleshoot issues with certain URLs
 - cached files can be purged easily with ‘rm’
 - use Spread to distribute and coordinate cache purging operations



The Next Step

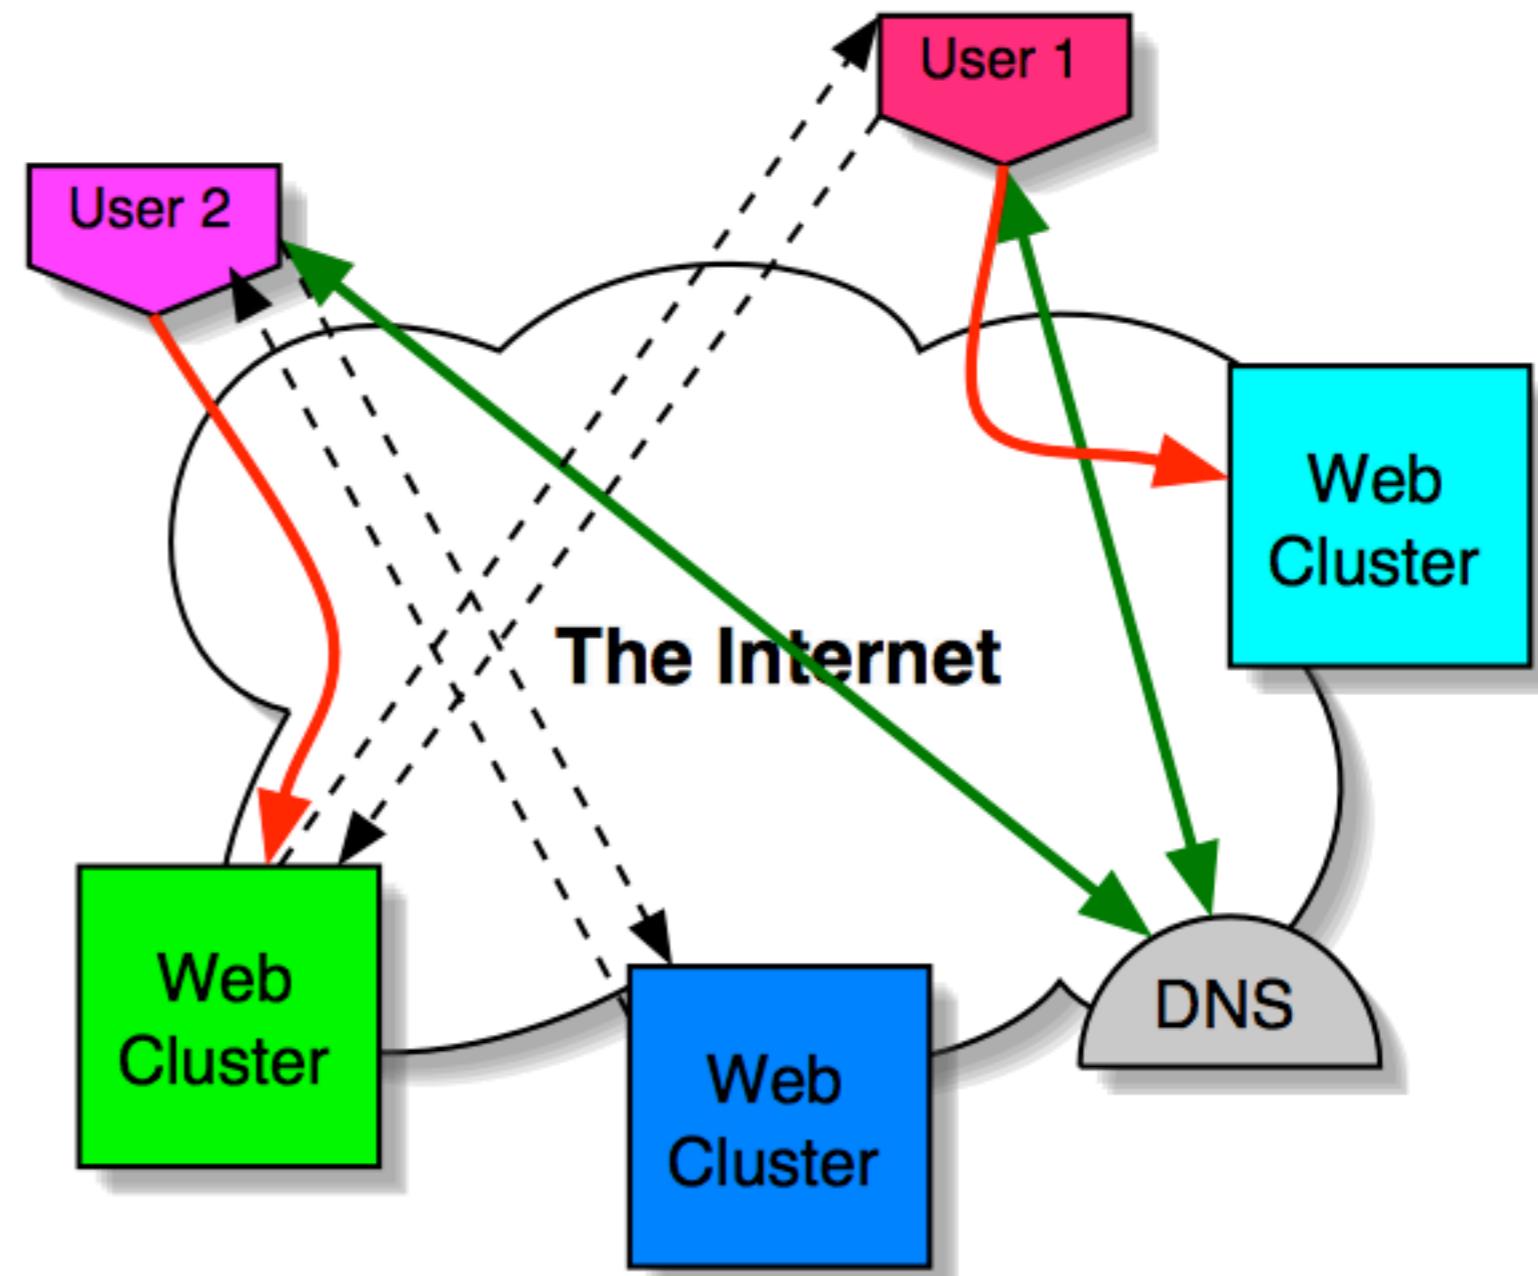
- ➊ We have three mini-clusters installed and configured, ready to throw bits by the trillions
- ➋ How people find them?
...DNS, obviously
- ➌ How do people find the closest cluster to them?
...clever DNS, [not so] obviously



Replica Location

server-side redirection

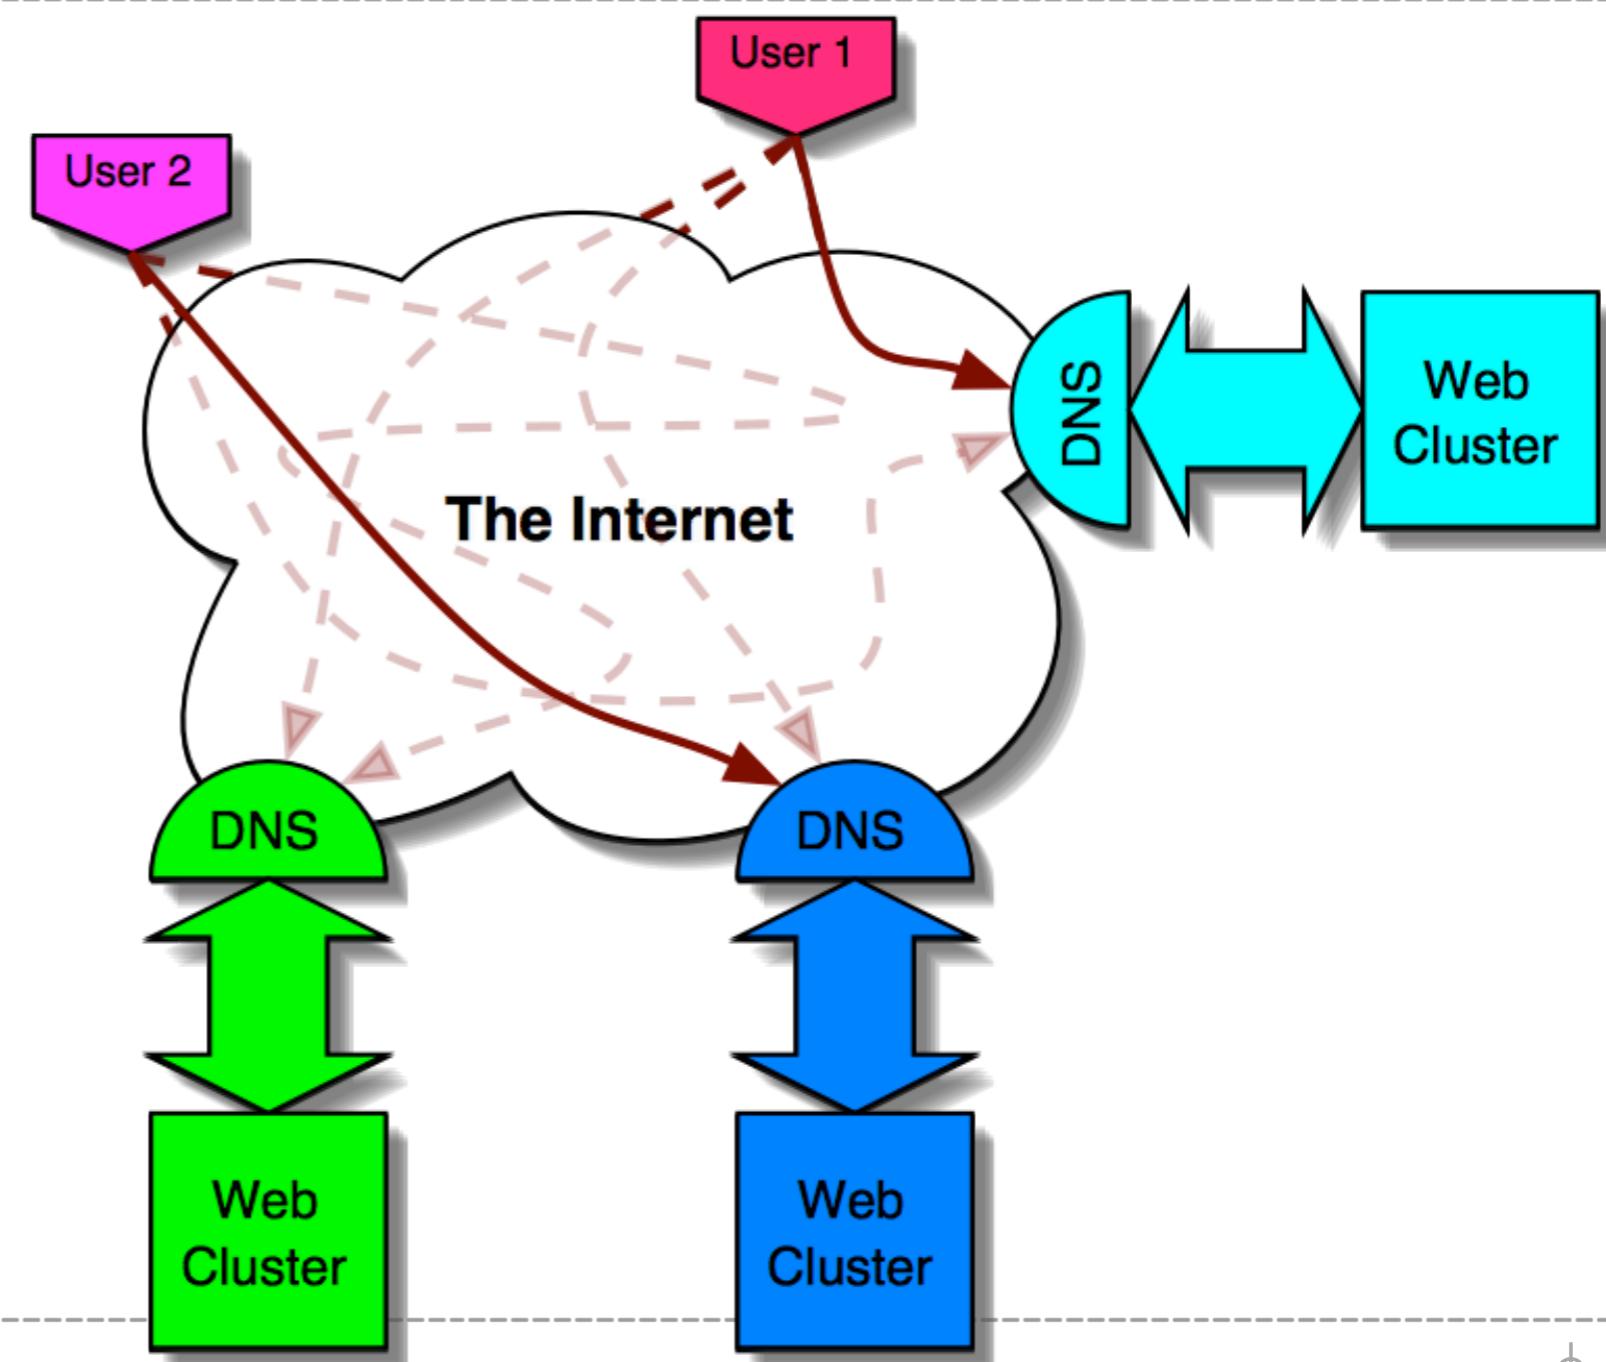
- Application is responsible
- Use HTTP redirects:
 - images-sj.example.com
 - images-nj.example.com
 - images-de.example.com



Replica Location

Proximity-based DNS

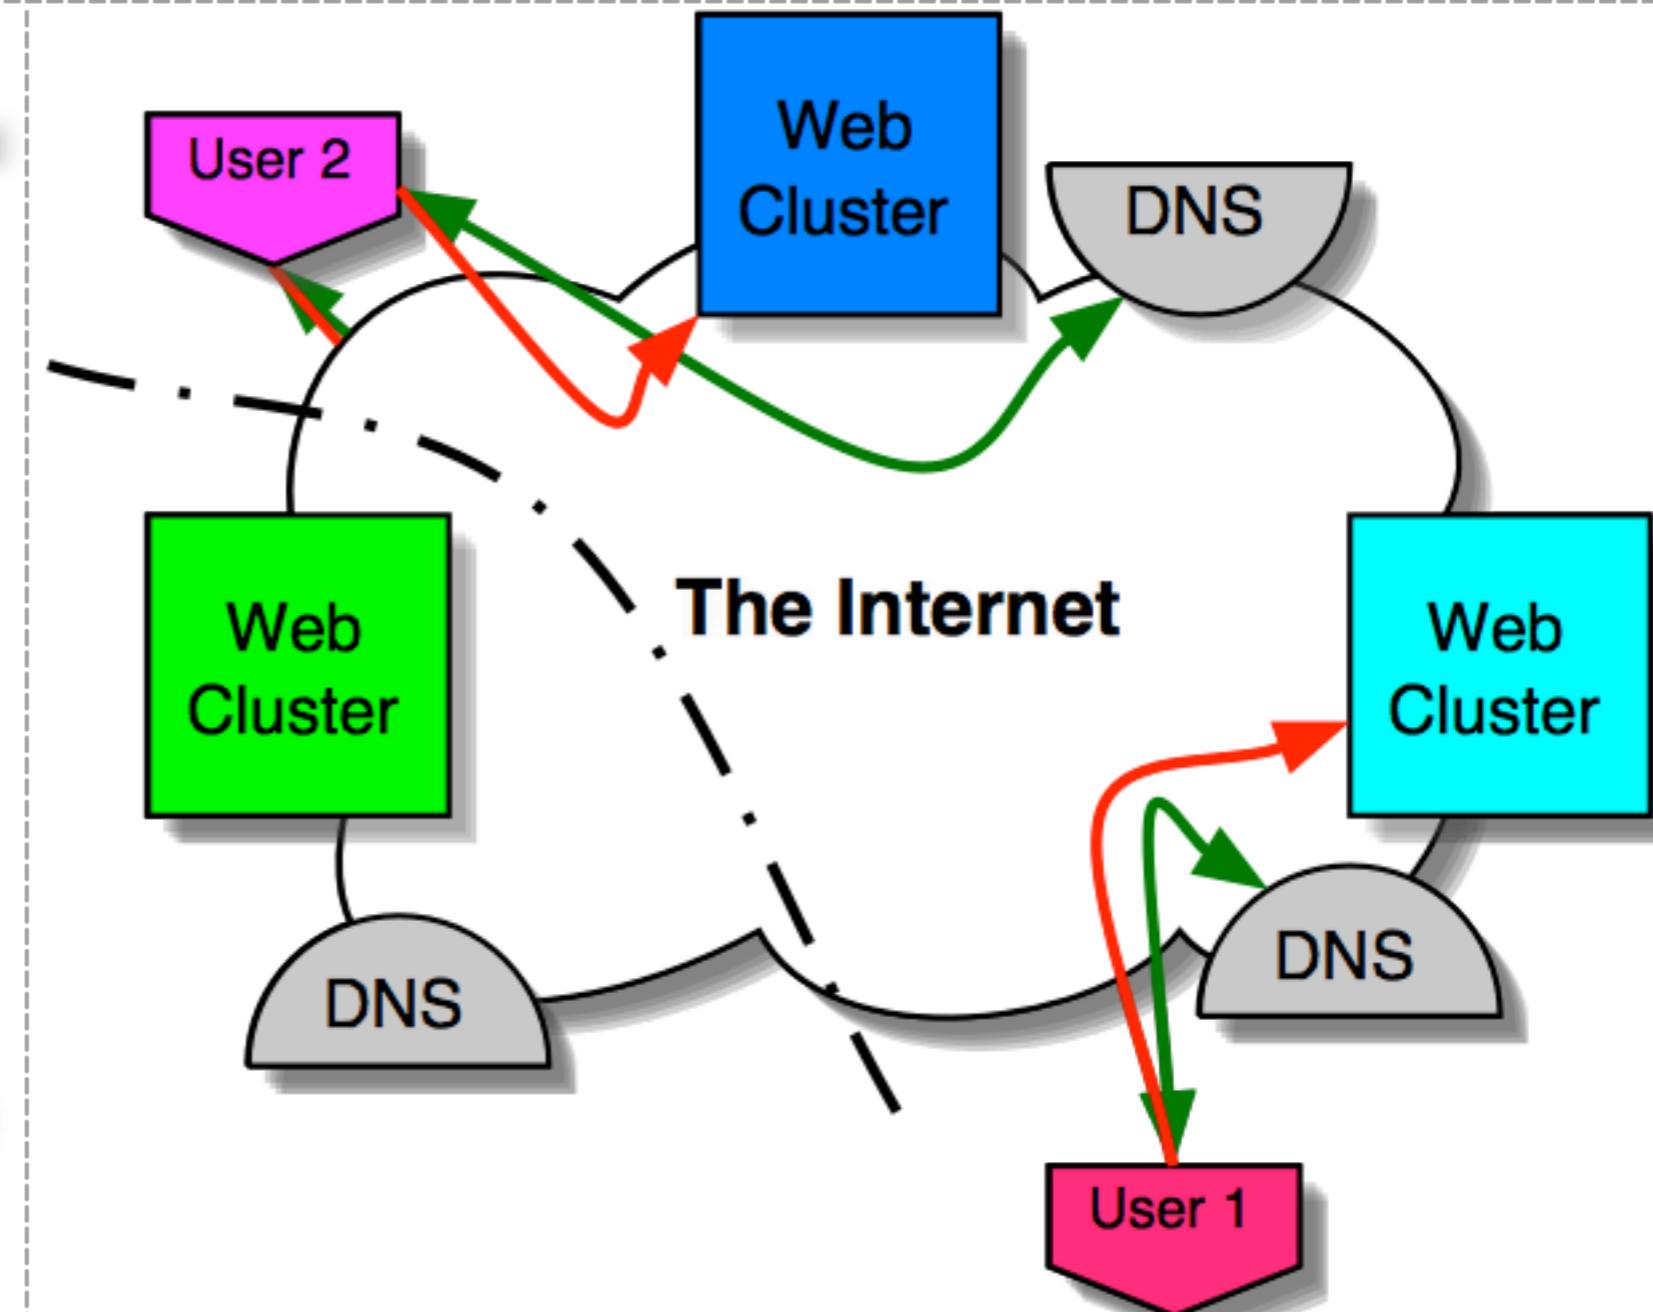
- DNS Provides Convergence
- Can take hours/days
- DNS servers near clusters:
 - ns-sj.example.com
 - ns-nj.example.com
 - ns-de.example.com



Replica Location

DNS Shared IP (a.k.a. AnyCast)

- DNS servers near clusters:
 - ns-sj.example.com
 - ns-nj.example.com
 - ns-de.example.com
- All DNS servers have the same IP
- Network block is announced from all sites via BGP
- Routing protocols provide immediate convergence



Web Cluster Logging



The Setup and The Goal

- Cluster of web servers
 - Apache
 - thttpd
- Logs are vital
 - must be stored in more than one place
- Real-time assessments
 - hit rates
 - load balancing
 - HTTP response code rates



Traditional Configuration

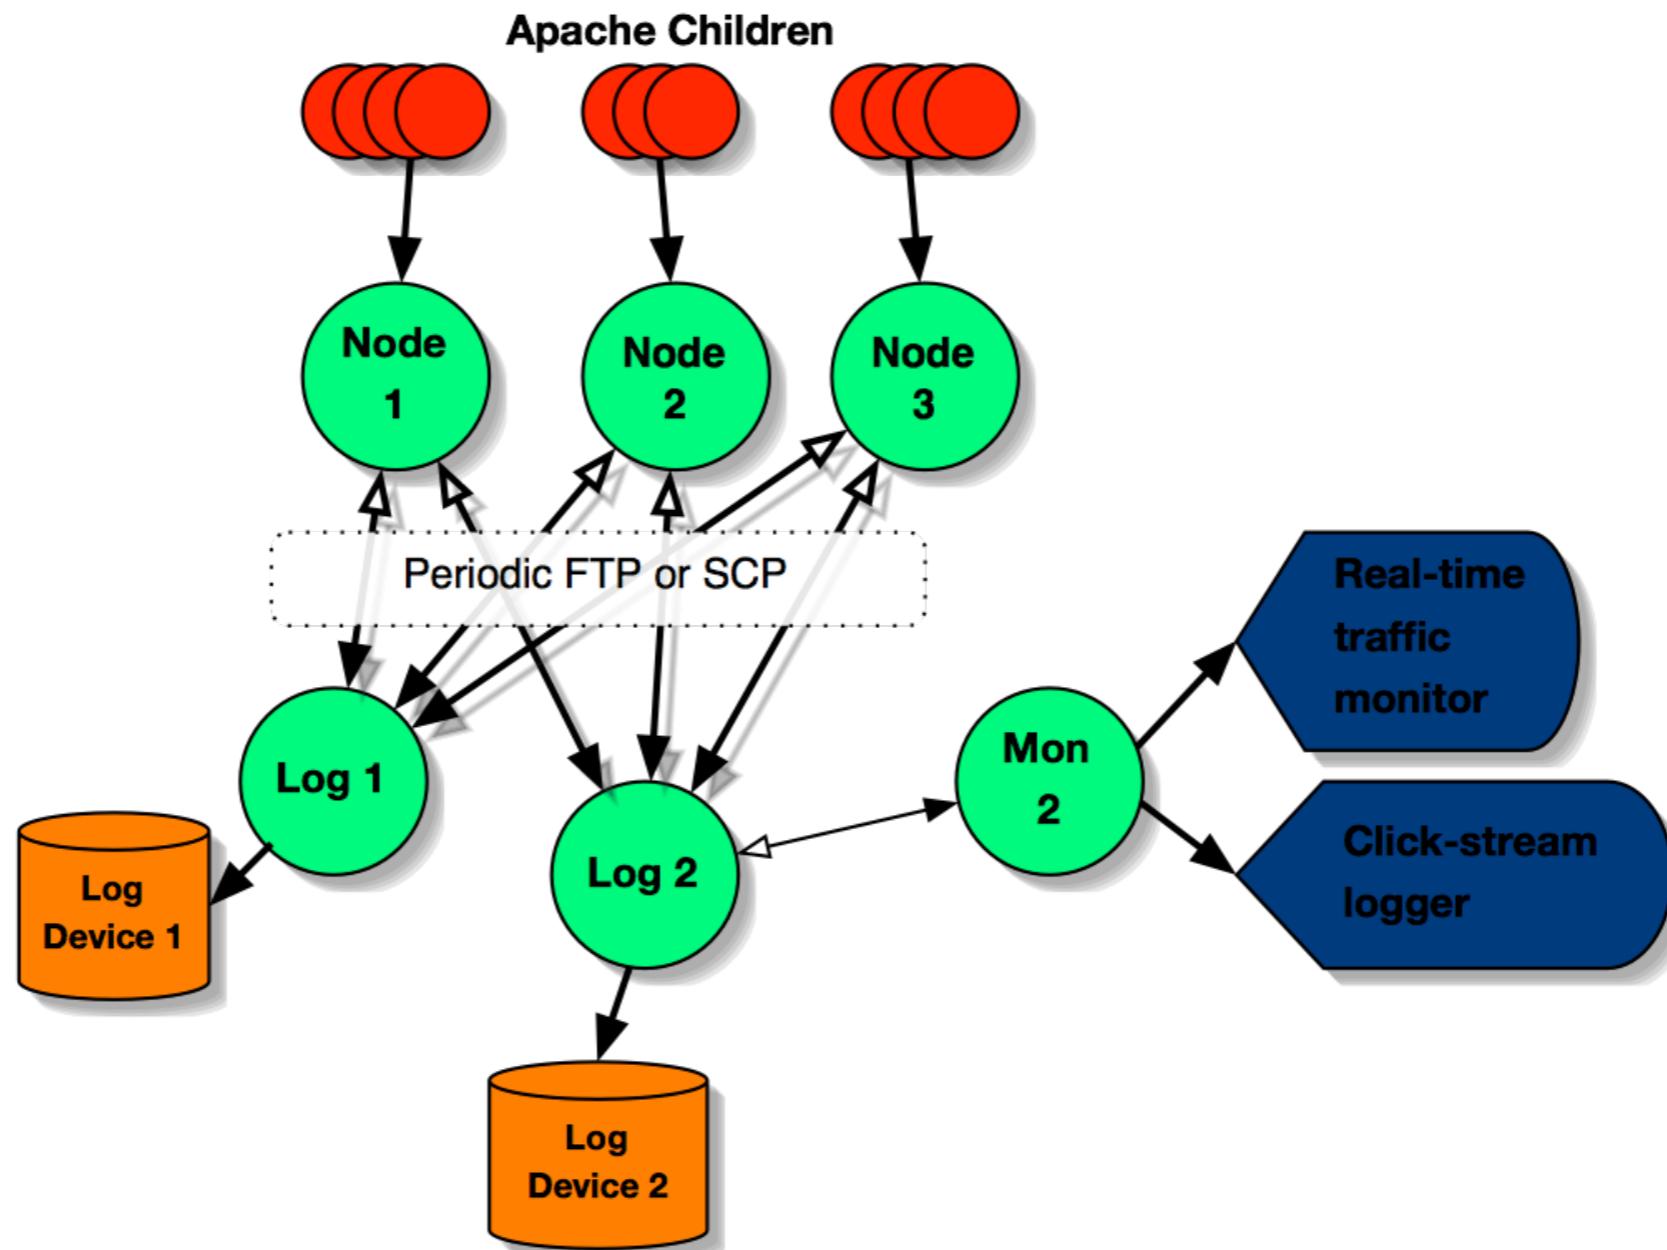
Local Logging, Post-process Aggregation

- Log written locally on web servers
 - space must be allocated
- Consolidation happens periodically
 - crashes will result in missing data
 - aggregators must preserve chronology (expensive)
 - real-time metrics cannot be calculated
- Monitor(s) must run against log server
 - monitors must tail log files
 - requires resources on the logging server



Traditional Configuration

Local Logging, Post-process Aggregation



TCP/IP or UDP/IP Logging

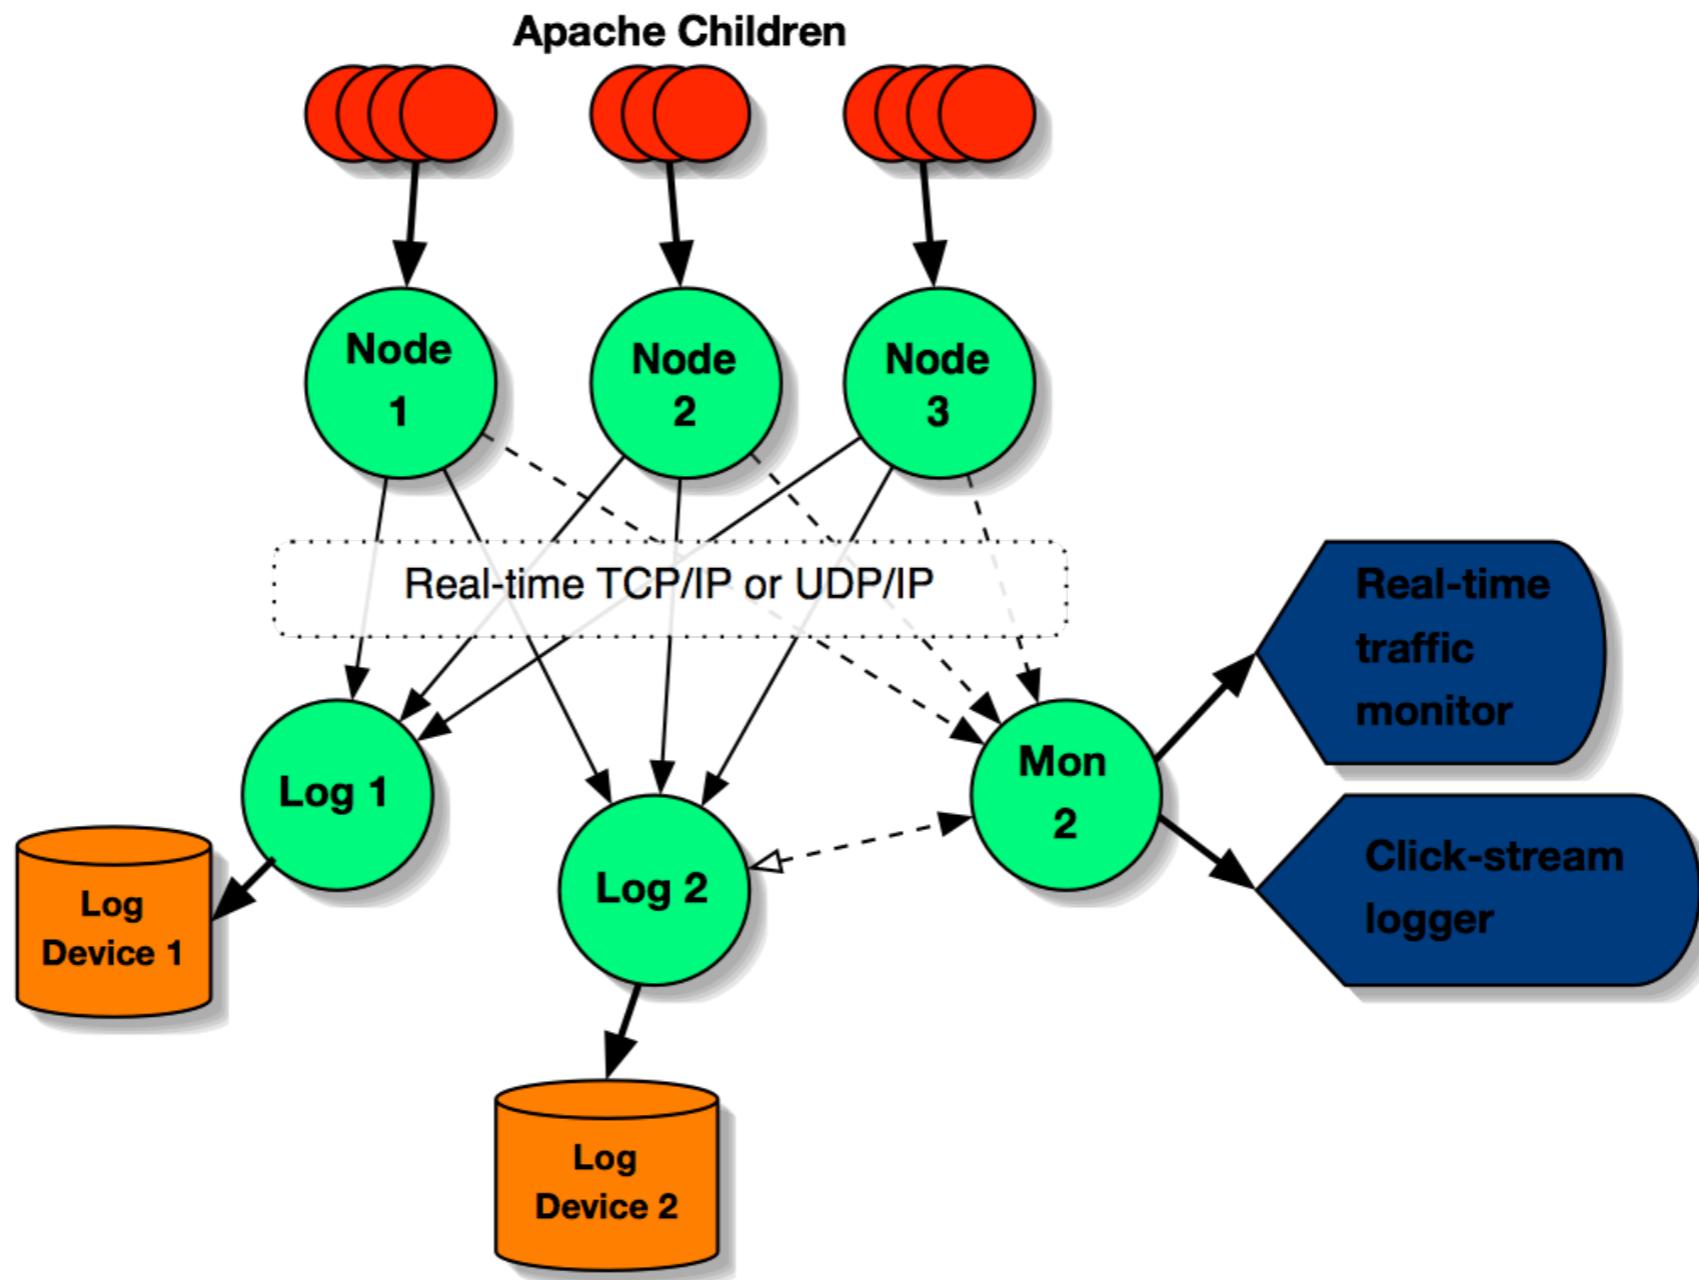
Syslog, Syslog-NG

- Logs are written directly to logging server(s)
 - UDP is unreliable and thus not useful
 - TCP is a point-to-point protocol
 - Two logging servers means all info gets sent twice.
 - Add a monitor and that's three times!
- Real-time metrics can now be collected
 - monitors must still be run against log servers
(or the publishers must be reconfigured)



TCP/IP or UDP/IP Logging

Syslog, Syslog-NG



Passive Network Logging

sniffers

- **Requires no modification of architecture**
 - Add/remove publishers (web servers) on-the-fly
- **Drops Logs!**
 - When tested head-to-head with traditional logging mechanisms we see loss
 - “Missing” logs are unacceptable
- **Clean the white dust off your upper lip and choose another implementation**



Reliable Multicast Logging

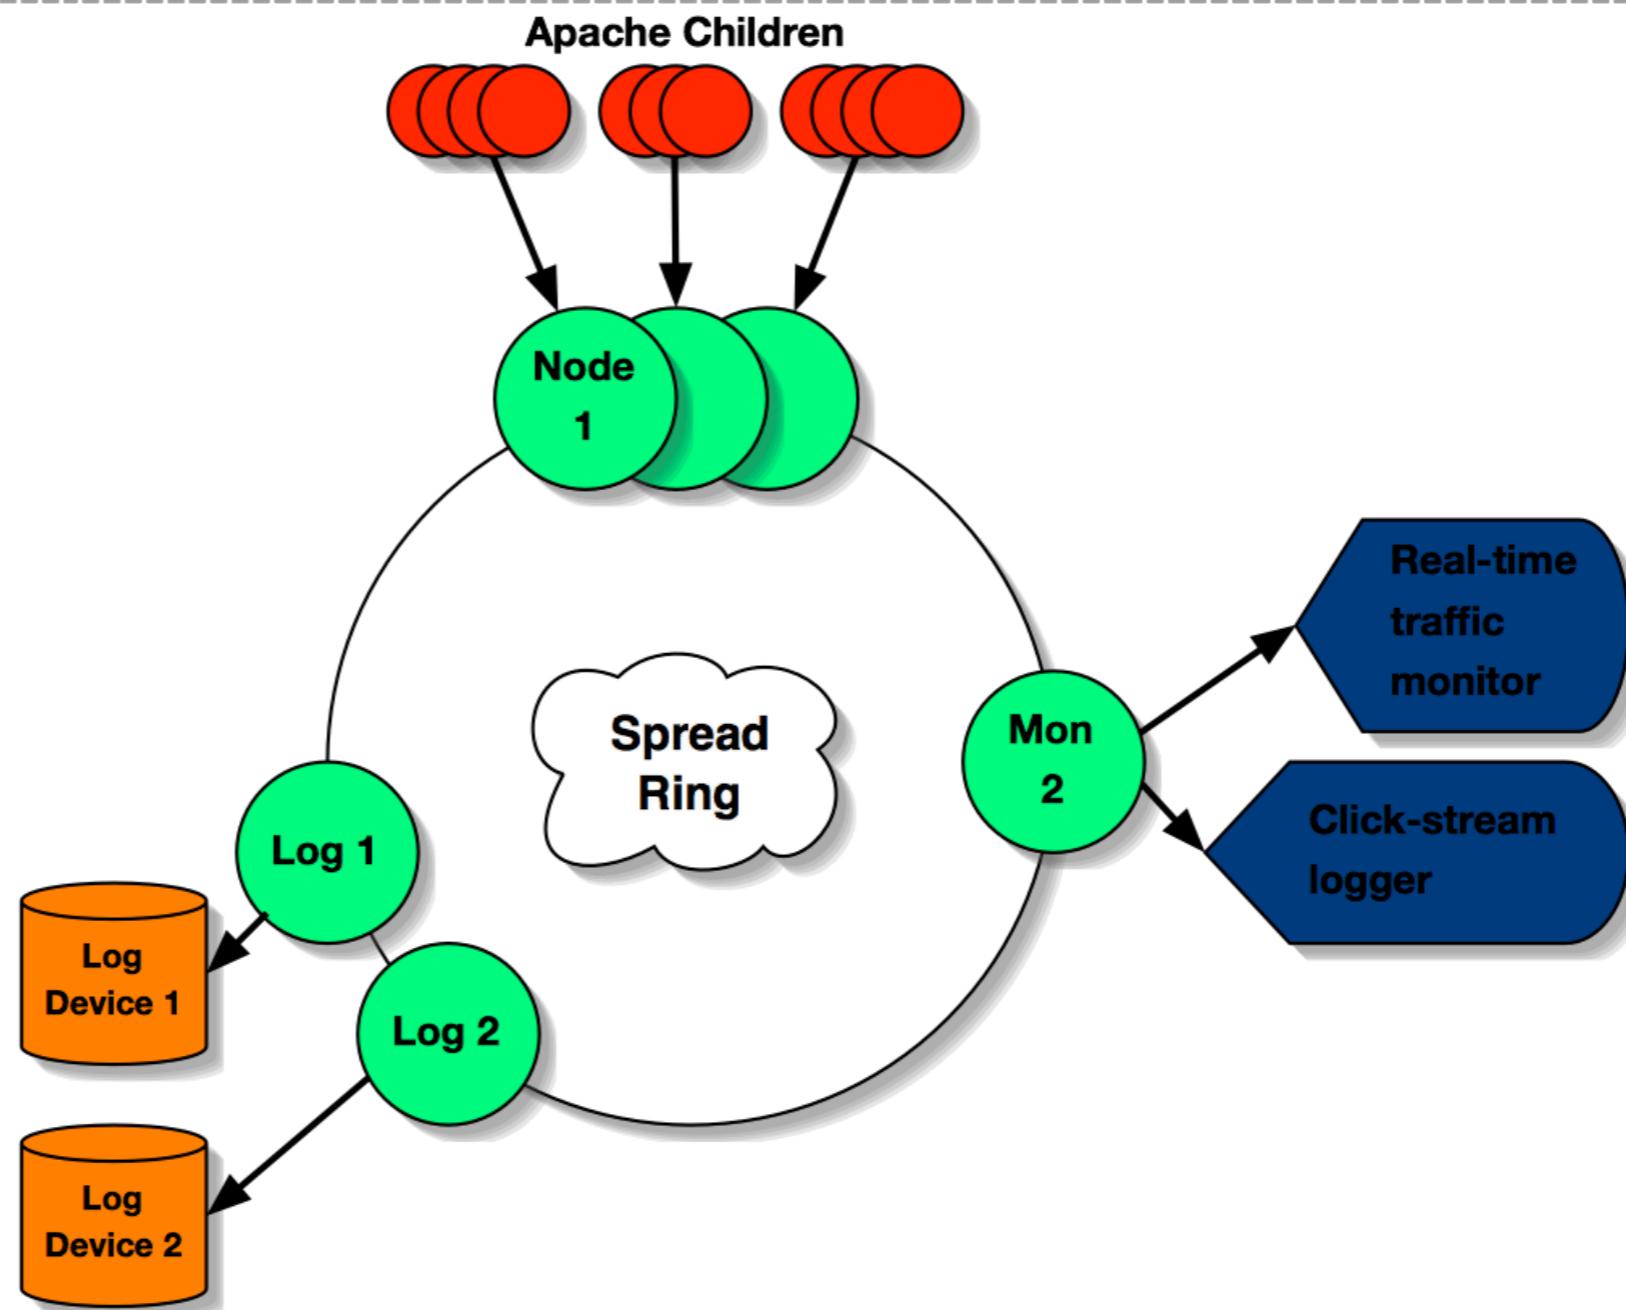
mod_log_spread

- **Flexible Operations**
 - Add/remove publishers (web servers) on-the-fly
 - Add/remove subscribers (loggers/monitors) on-the-fly
- **Reliable Multicast (based on Spread)**
 - Multiple subscribers don't incur addition network overhead
- **Individual real-time passive and active monitors**
 - Monitors can be “attached” without resource consumption concerns
 - Passive monitors that draw graphs, assess trends, detect failures
 - Active monitors that feed metrics back into a production system
 - Who's online
 - Real-time page access metrics



Reliable Multicast Logging

mod_log_spread





mod_log_spread

“The Publisher”

mod_log_spread is really a patch to mod_log_config

Like pipes in mod_log_config:

“| /path/to/rotatelogs filename 3600”

mls adds a Spread group destination:

\$groupname

```
LogFormat "%h %l %u %t \"%r\" %>s %b \"%{Referer}i\" \"%{User-Agent}i\" %T" combined  
CustomLog $example combined
```





spreadlogd

“The Subscriber”

spreadlogd writes logs...

BufferSize = 65536

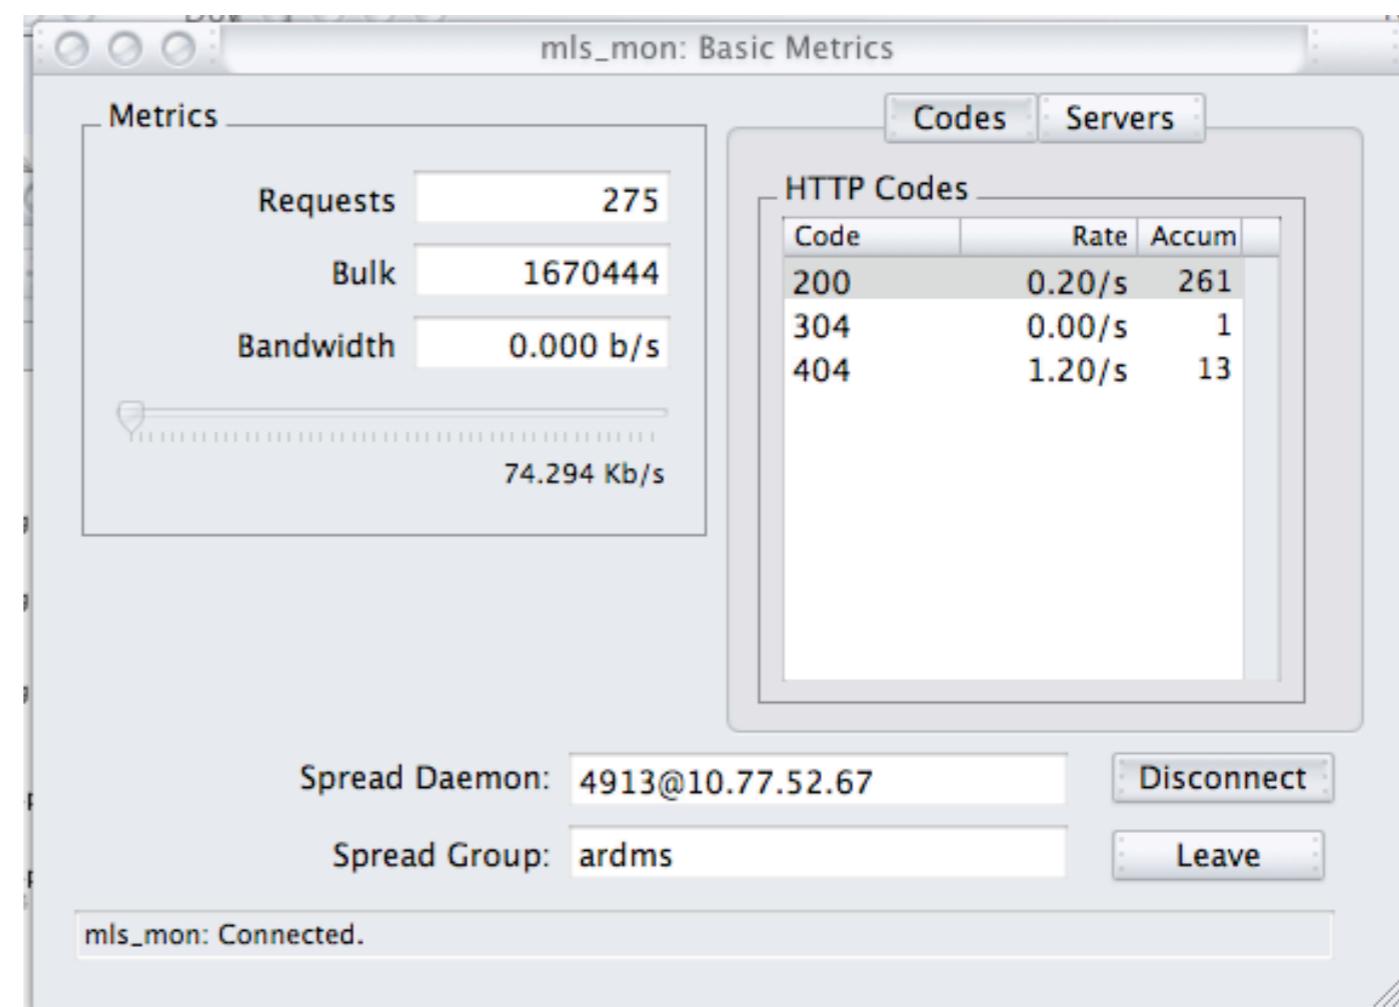
```
Spread {
    Port = 4803
    Log {
        RewriteTimestamp = CommonLogFormat
        Group = "example"
        File = /data/logs/apache/www.example.com/combined_log
    }
}
```

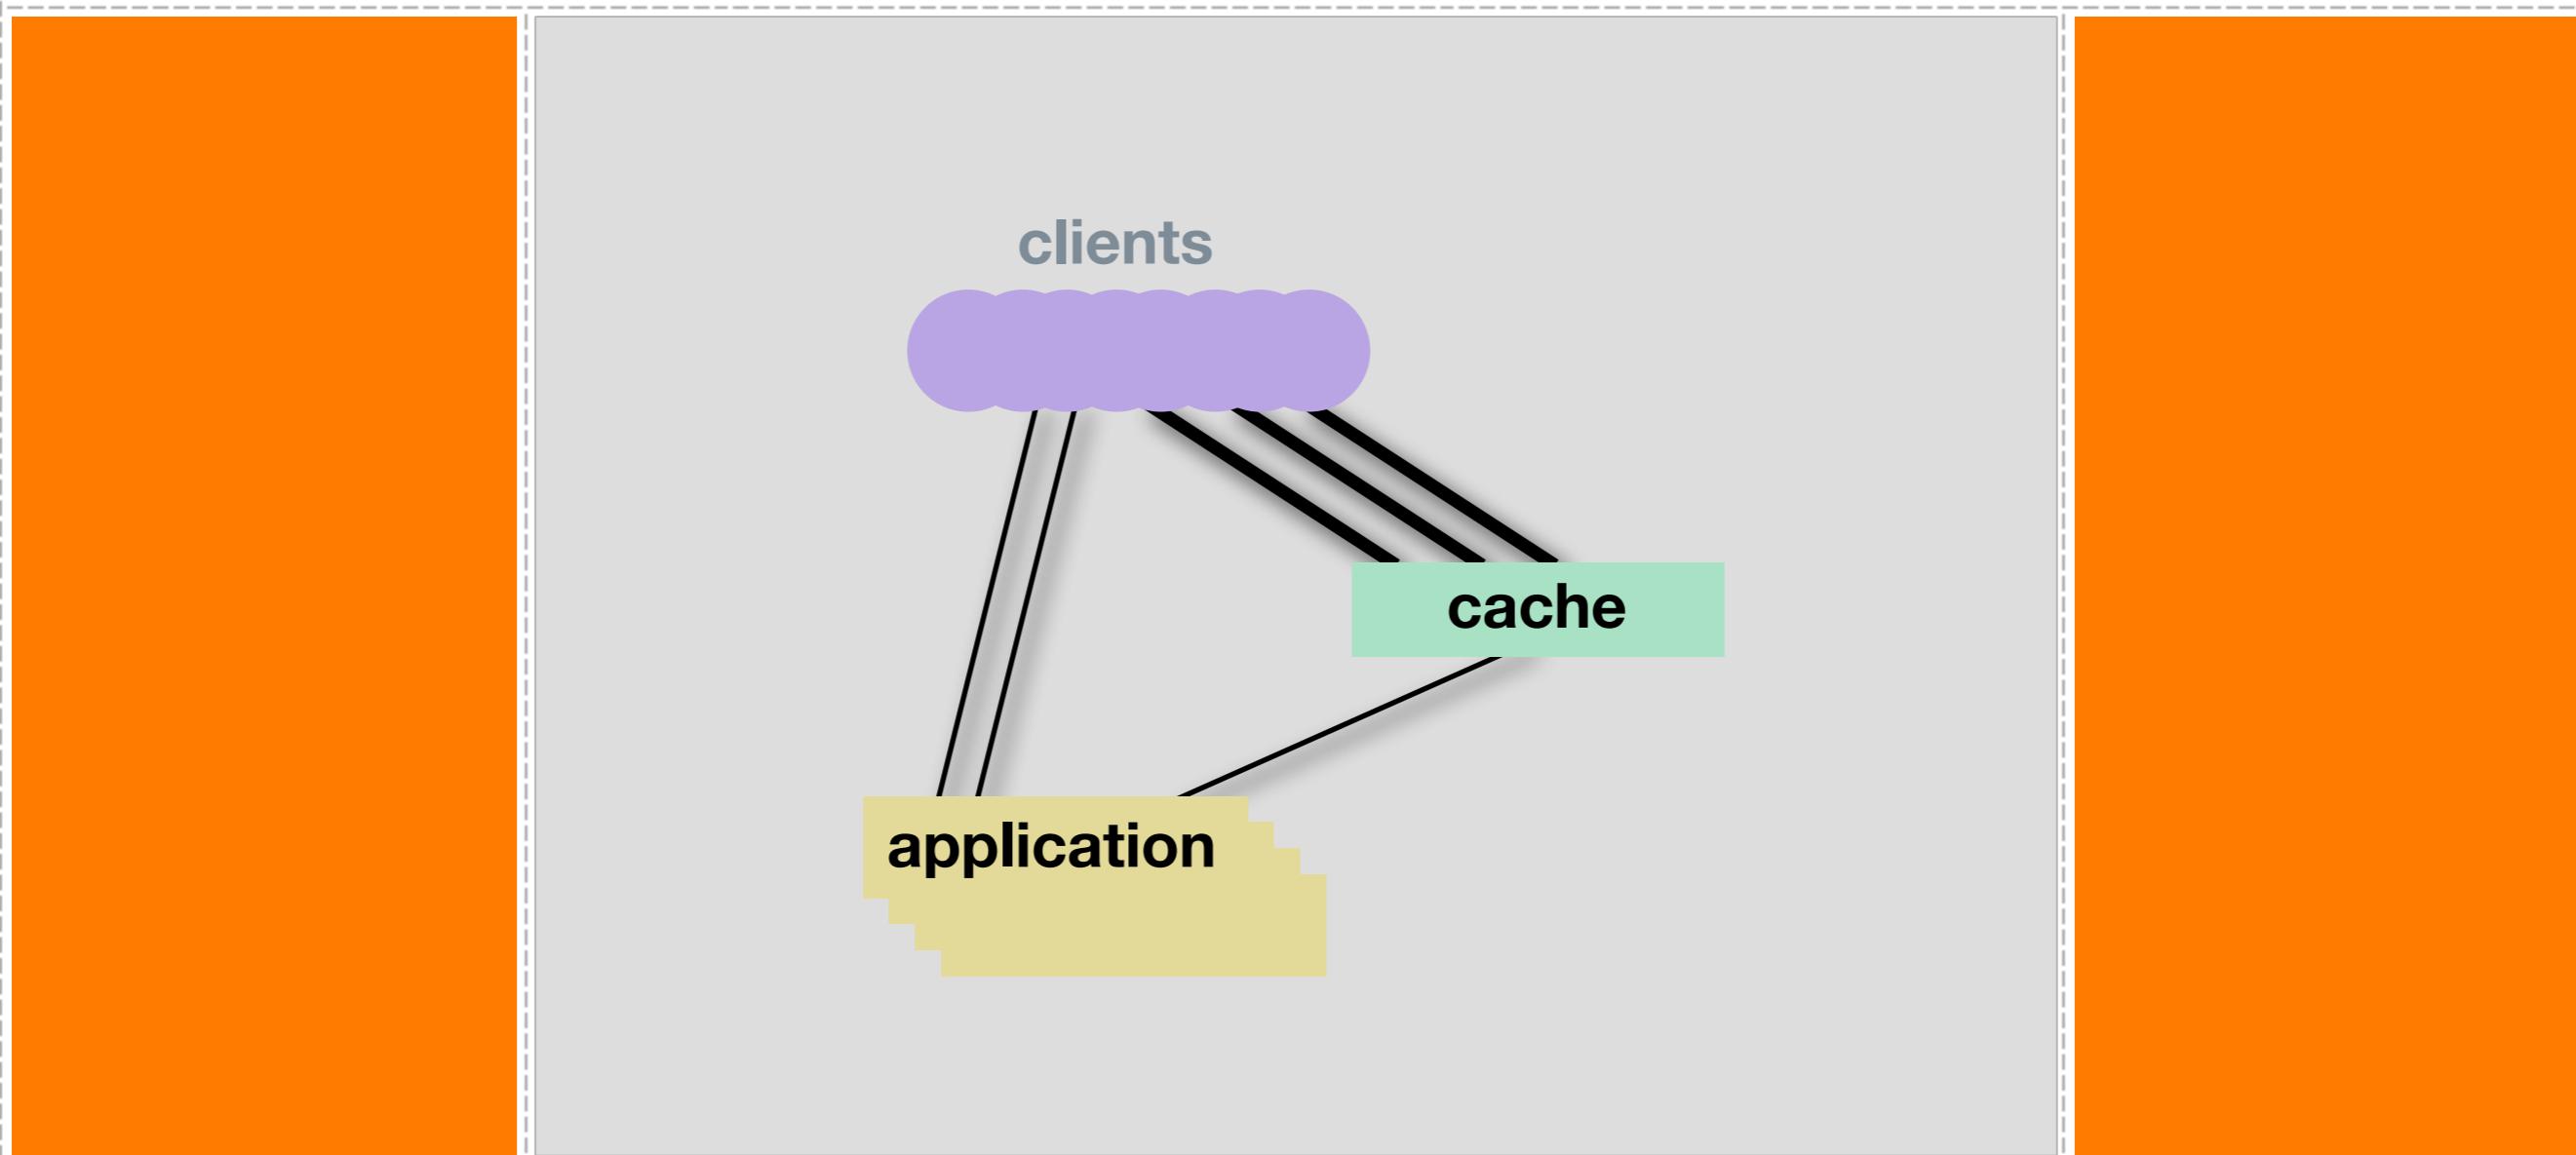


Other Real-Time Tools

“Other Subscribers”

- **ApacheTop**
- C++ “top”-style real-time hit display
- **mls_mon**
- graphical hit rates by server and by code





Caching Architectures





What is a cache?

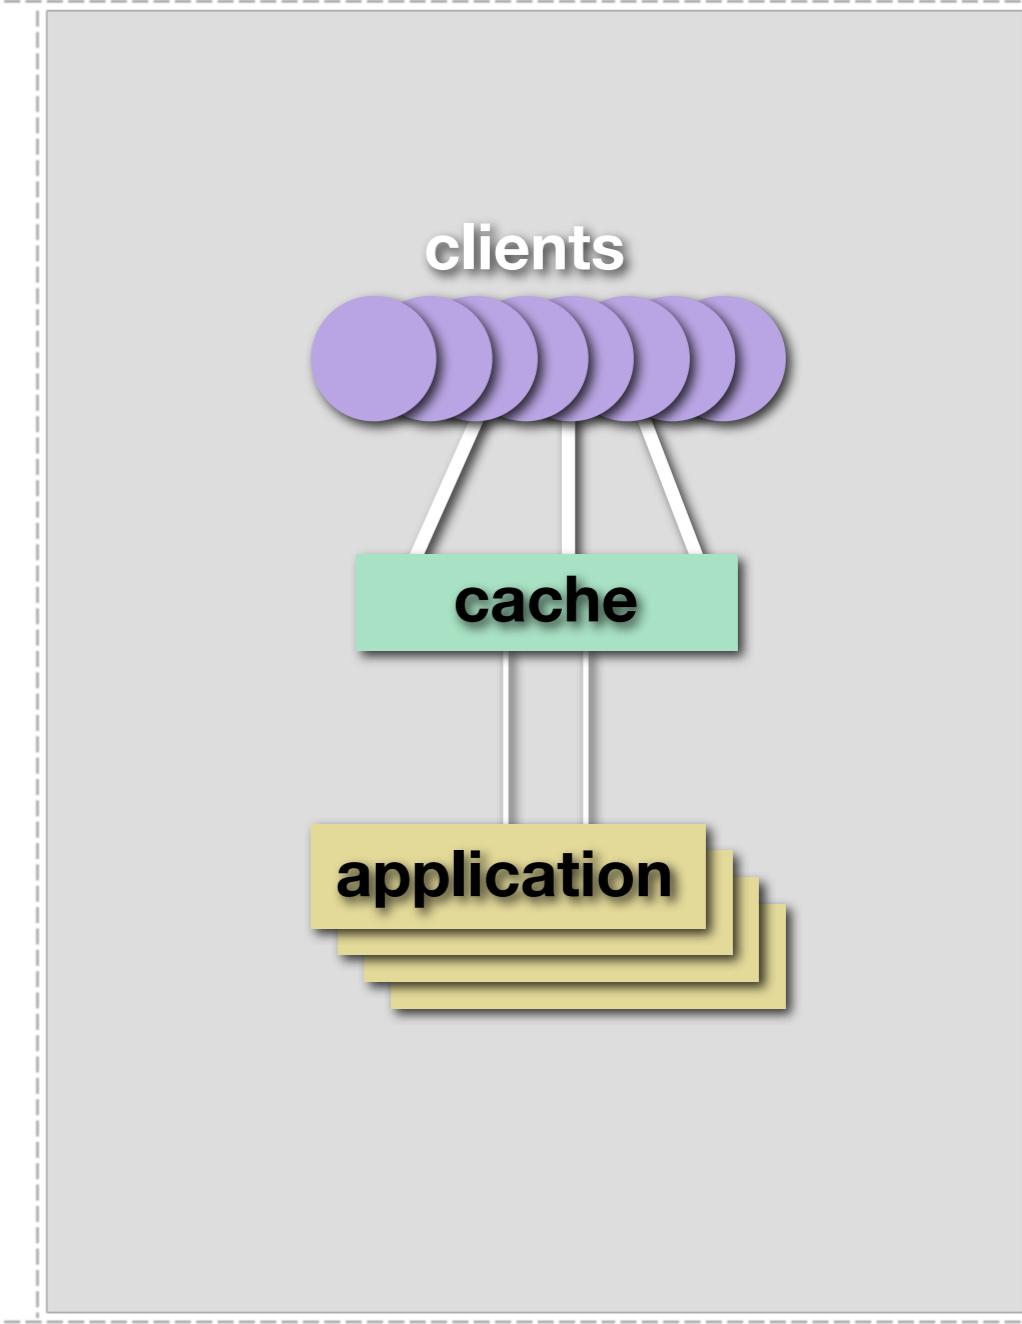
Cache:

A small fast memory holding recently accessed data, designed to speed up subsequent access to the same data. Most often applied to processor-memory access but also used for a local copy of data accessible over a network etc.



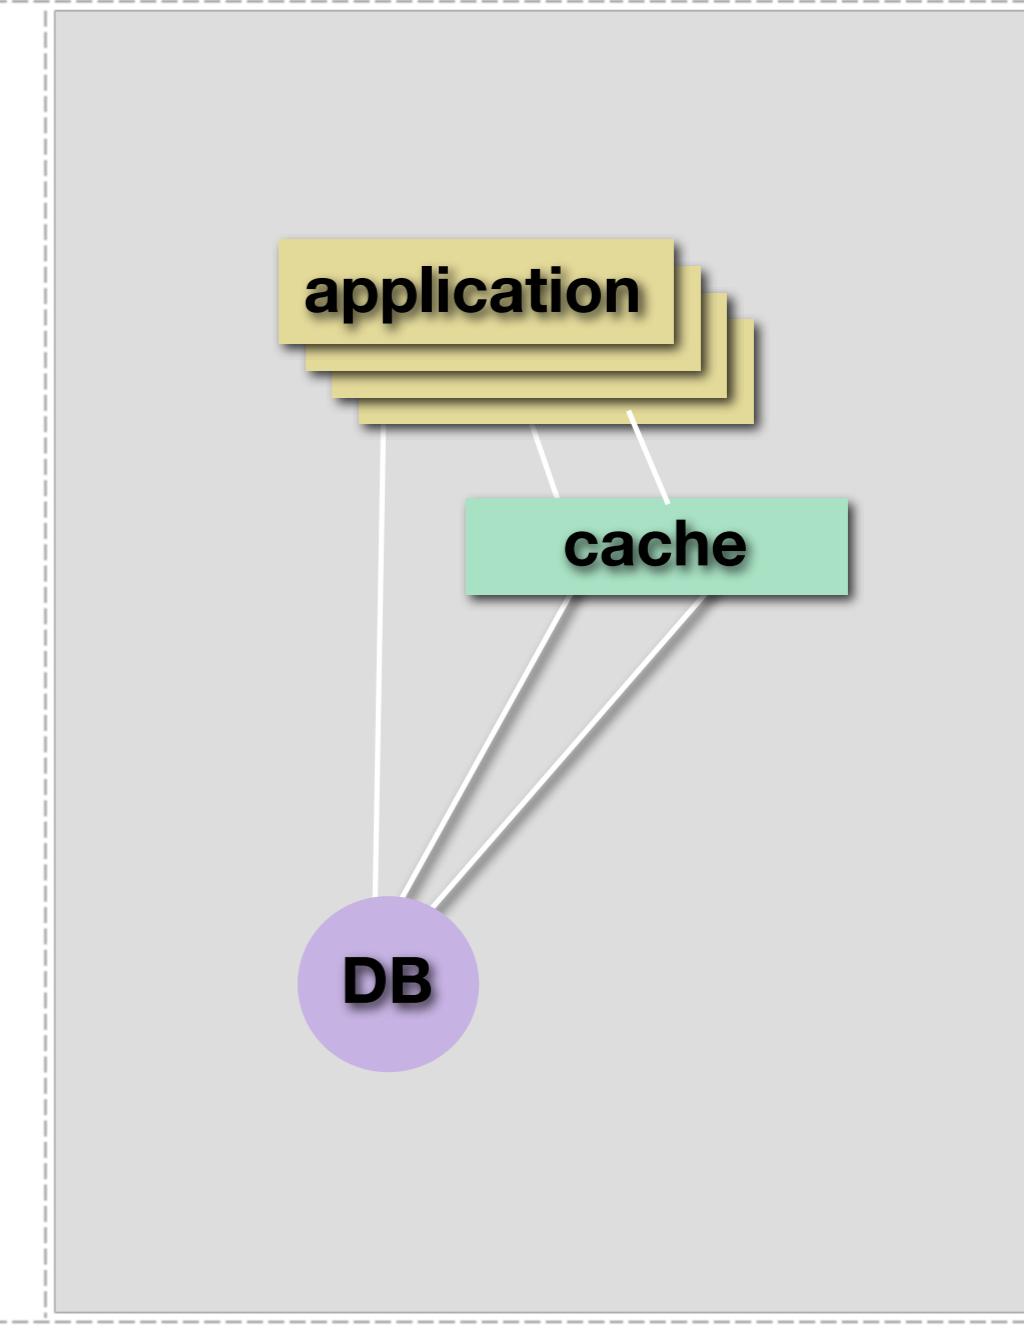
The Layered Cache

- Exists above/in-front
- Knows little or nothing about what's underneath
- Works fabulously for static data (like images)



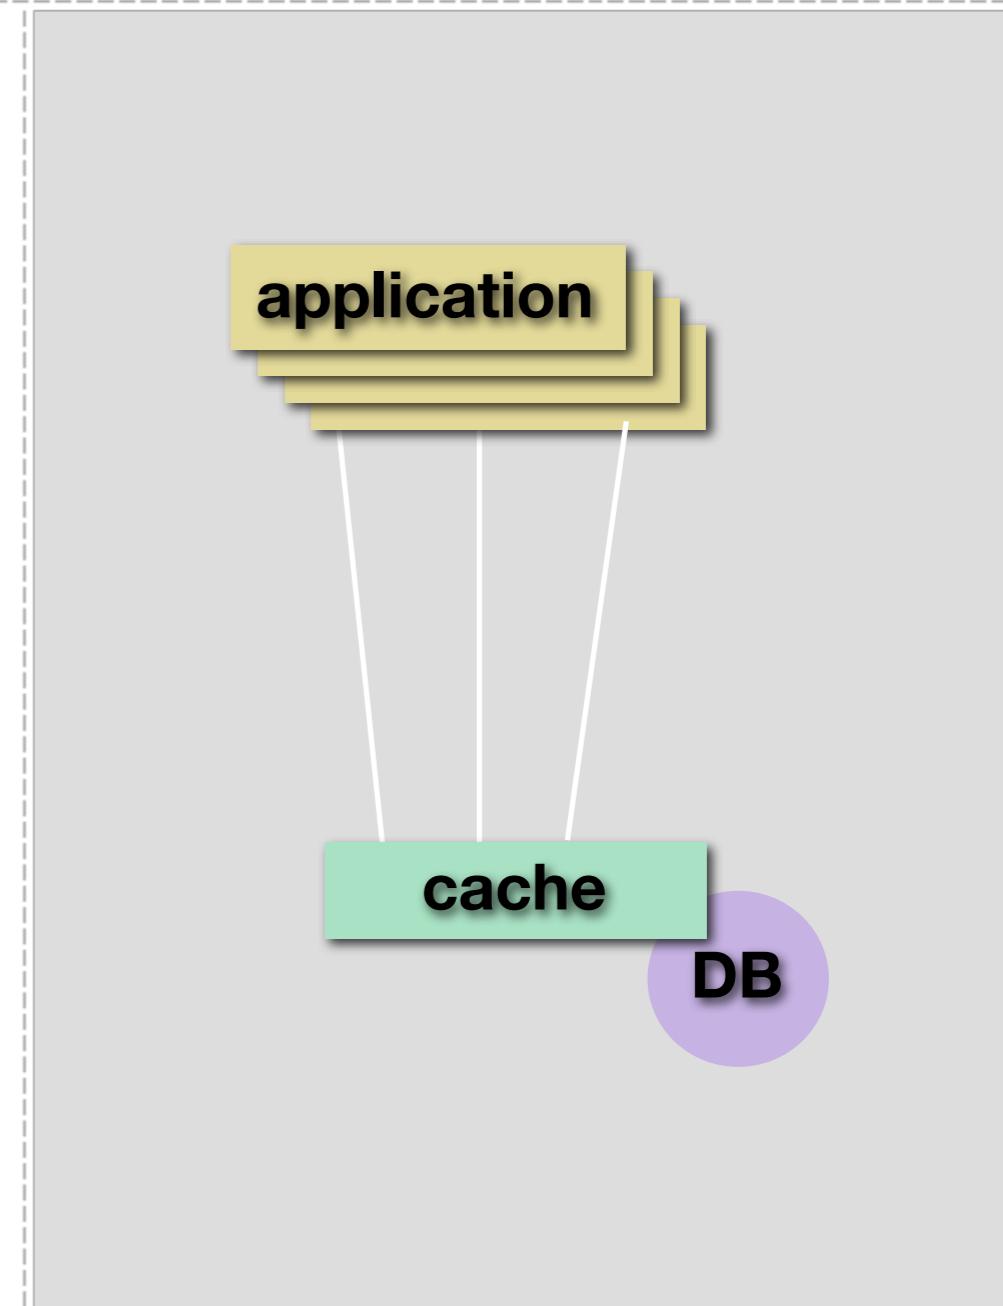
The Integrated Cache

- Exists in the application.
- Knows the data and how the application uses it.
- Works well for data that doesn't change rapidly but is relatively expensive to query.



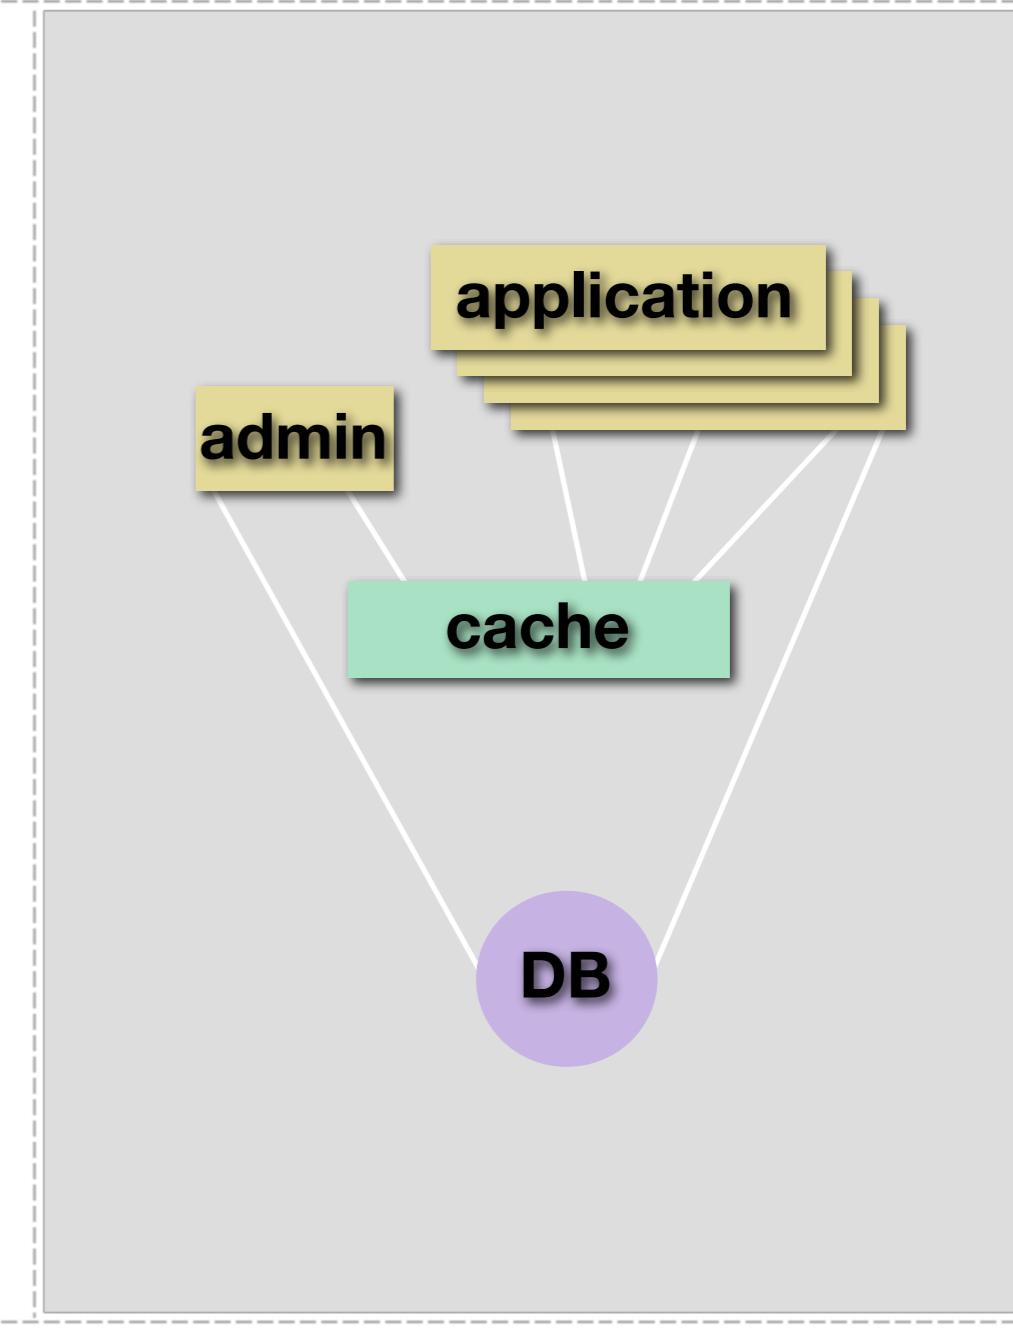
The Data Cache

- Exists in the data store
- Knows the data, the queries and how the data has changed.
- Works well always
 - called computational reuse
 - oldest trick in the book
- MySQL 4 has this
they call it a Query Cache



Write-thru Cache

- Occurs at update location
- Knows the data, the app, the queries and how the data has changed.
- Works well for administrative updates
 - many WebLogs work this way
- Can be very adaptive and flexible



A “Real-World” Example

- News site
 - News items are stored in Oracle
 - User Preferences are stored in Oracle
- Hundreds of different sections
 - Each with thousands of different articles
- Pages:
 - 1000+ hits/second
 - shows personalized user info on **EVERY** page
 - front page shows top N_F articles for forum F (limit 10)





The Approach

- Oracle is fast enough
 - why abuse Oracle for this purposes?
 - surely there are better things for Oracle to be doing

- Updates are controlled
 - updates to news items only happen from a publisher
 - news update:read ratio is minuscule
 - user preferences are only ever updated by the user



Articles

- Article publishing
 - sticks news items in Oracle
- The straight forward way
 - <http://news.example.com/news/article.php?id=12345>
 - page pulls user prefs from cookie
(or bounces off a cookie populator)
 - page pulls news item from database
- I hate query strings
 - I like: <http://news.example.com/news/items/12345.html>

```
RewriteRule ^/news/items/([^\"]*)\.html$ /www/docs/news/article.php?id=$1 [L]
```



Articles Cached

- We pull the item that is likely to never change
 - cheaper if the page just hard coded the news item
 - writing the news article out into a PHP page is a hassle
 - ... or is it?
- Have the straight forward page cache it
 - /news/article.php writes /news/items/12345.html as a PHP page that still expands personal info from cookie, but has the news item content statically included as HTML.

```
RewriteCond %{REQUEST_FILENAME} ^/news/items/([^\"]*)\.html
RewriteCond %{REQUEST_FILENAME} !-f
RewriteRule ^/news/items/([^\"]*)\.html$ /www/docs/news/article.php?id=$1 [L]
```



Articles Cache Invalidation

- Run a cache invalidator on each web server
 - connects to Spread as a subscriber
 - accepts /www/docs/news/items/####.html deletion requests
 - accepts full purge requests
- Article publishing
 - stash item #### in Oracle (insert or update)
 - publish through Spread an invalidation of ####
- Changing the look of the article pages
 - change article.php to have the desired effect
(and write the appropriate php cache pages)
 - publish through Spread a full purge



The Result

- All news item pages require zero DB requests
- the business can now make your life difficult by requesting new crap on these pages that can't be so easily cached
- Far fewer database connections required
 - all databases appreciate that (Oracle, MySQL, Postgres)
- Bottleneck is now Apache+mod_php
 - crazy fast with tools like APC
 - inherently scalable... just add more web servers
 - room for more application features





Tiered Architectures



Why Tier?

- **Dedicate resources to specific components**
 - Often a good approach to scaling systems up
 - Requiring single purpose components is a good way to lock into a big (expensive) architecture
- **Tiers make computer science problems easier**
 - Understand the trade off of solving hard problems vs. maintaining tiered solutions
- **Tiers add complexity and increase maintenance costs**
 - More components, more pieces, more moving parts...
More can (and does) go wrong.





Dedicated Resources

- **Classic Example:**
 - Apache on dedicated web servers
 - Database on dedicated machine
 - Why? it is easy to have 4 web server, hard to have 4 databases

- **Lock-in Example:**
 - Web application on several servers
 - Requires local session state and sticky sessions
 - Why? scaling down to 2 servers will still require a load balancer that can “stick” sessions.



Tiering to Compensate (for problems that are hard to solve)

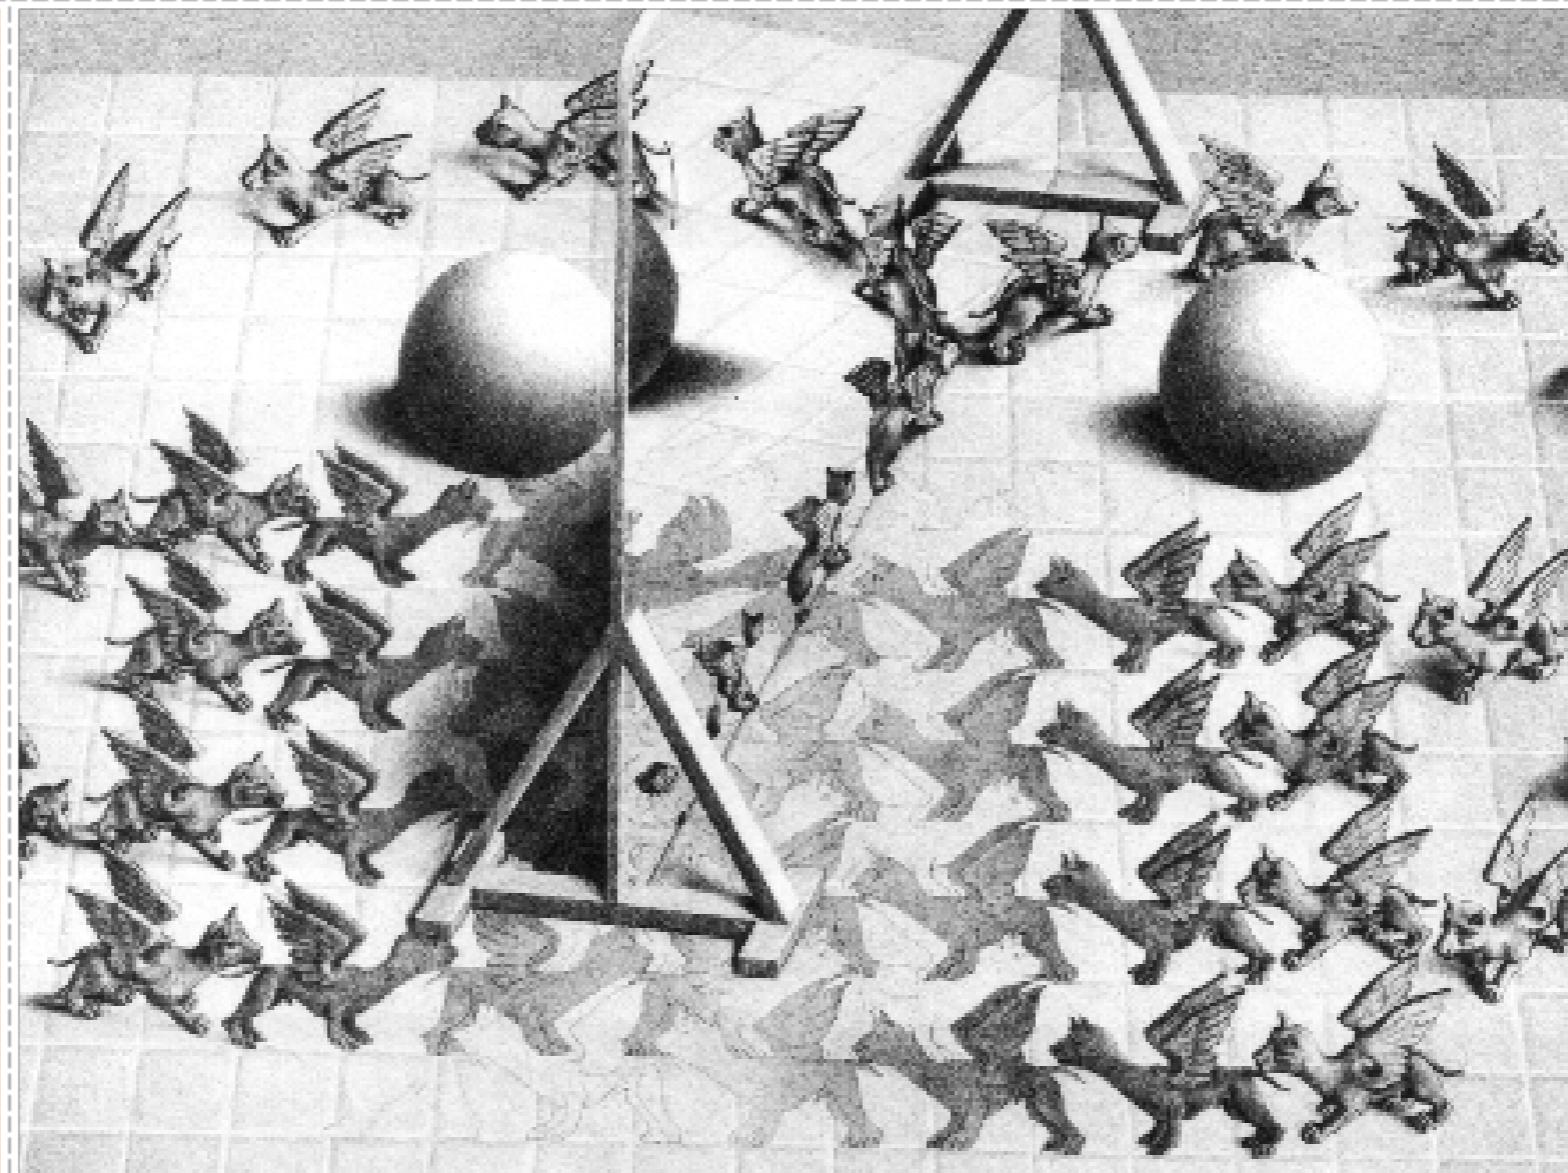
- Database Replication is hard
 - Anyone who tells you otherwise is lying or not telling you the whole story
- Session Replication is not so hard
 - use a technology like Splash!
 - offload responsibility to the client
- Tiers are expensive technically and financially
 - Some problems more difficult than tiers are expensive



Tiers are expensive

- Intrinsically difficult to scale down
- If it is a real production system...
 - Complete staging environment
 - Complete development environment





Effective Replication

Eliminating The Need To Tier

Types of Database Replication

Master-Slave

- **a data set has a master server**
- **changes to the data set are sent to slaves**
- **dml must be performed at the master**
- **read-only queries can be performed anywhere**
- **no challenging synchronization algorithms**



Types of Database Replication

Master-Master

- **data modification can be performed anywhere**
- **coordinating ACID and XA constraints is hard**
 - manage full transactions
 - view consistency
 - initial synchronization
- **synchronization algorithms**
 - 2-phase commit (2PC)
 - 3-phase commit (3PC)



Types of Database Replication

Multi-Master

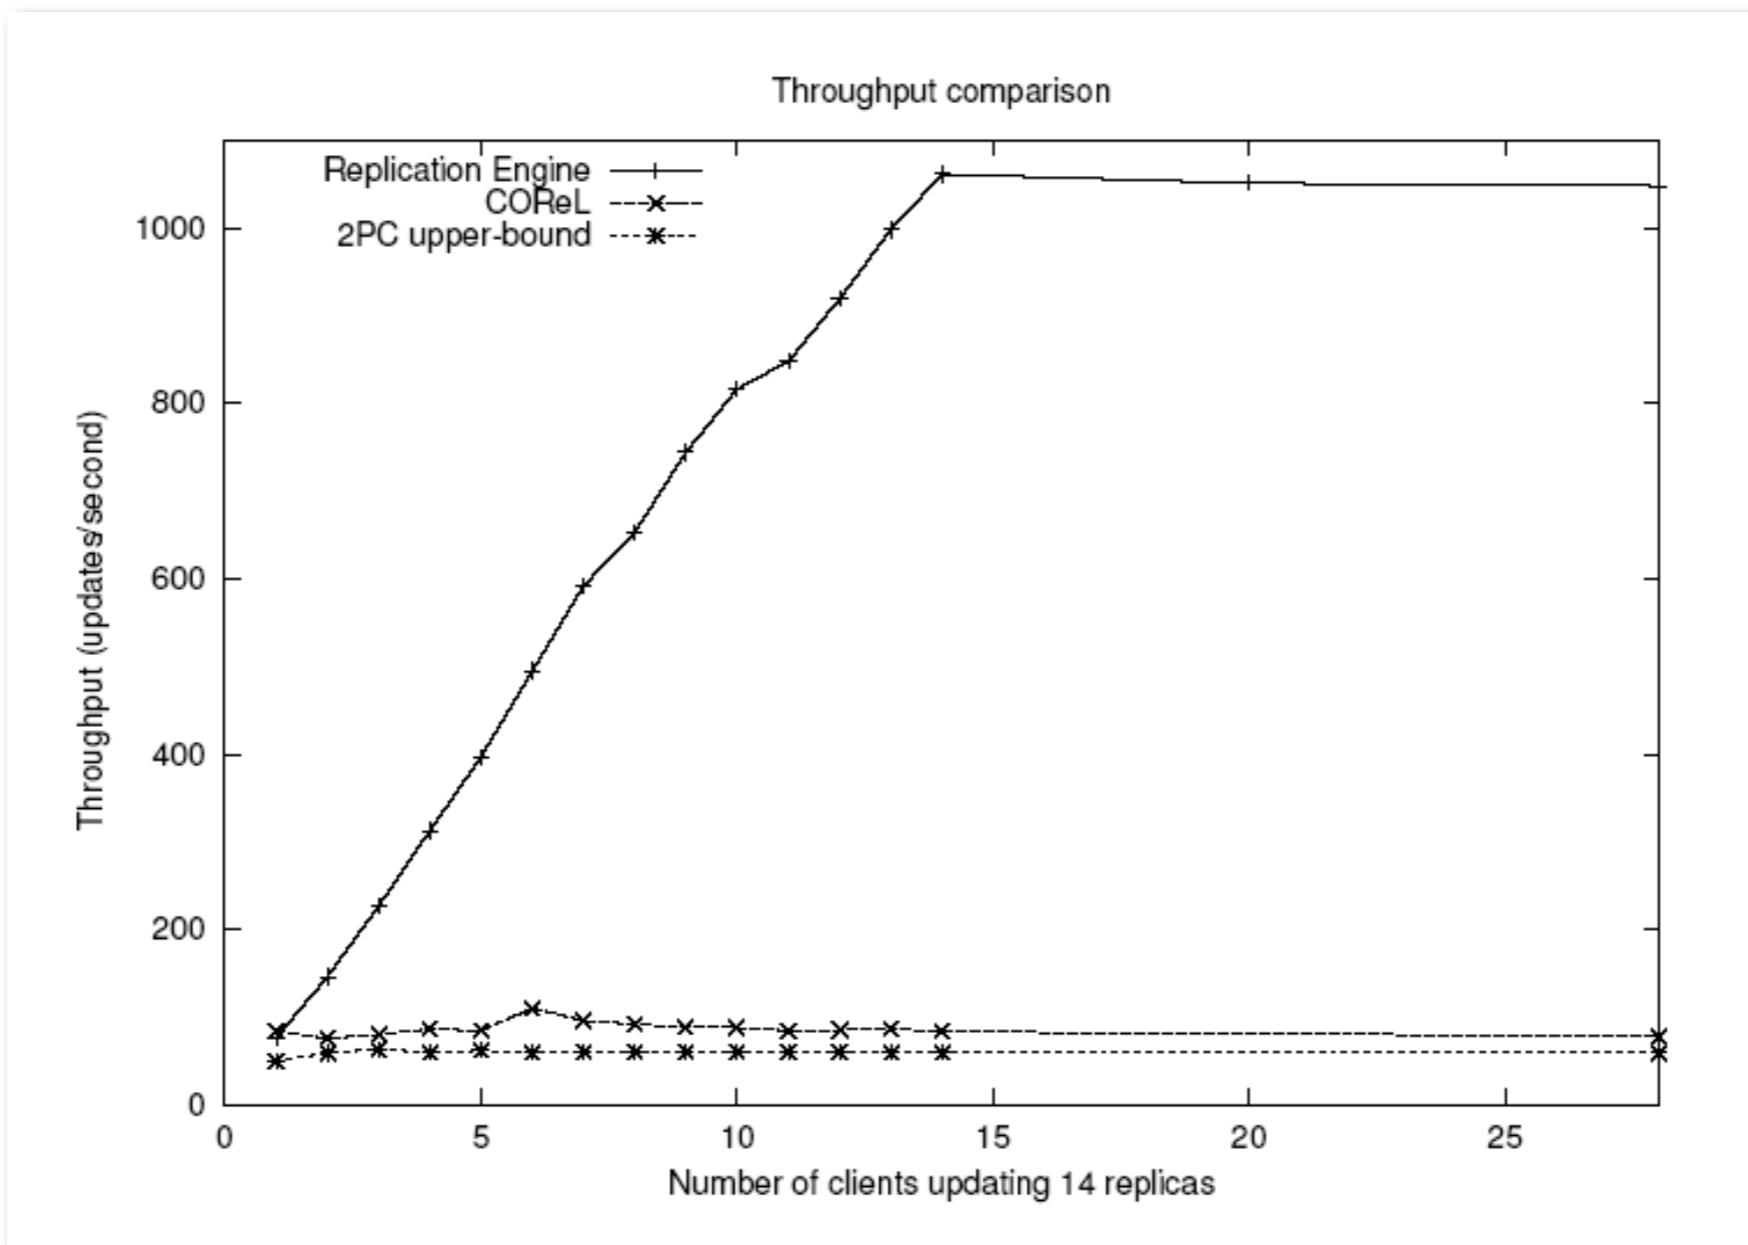
- **data modification can be performed anywhere**
- **coordinating ACID and XA constraints is hard**
 - manage full transactions
 - view consistency
 - initial synchronization
- **synchronization algorithms are complex**
 - 2PC and 3PC are unrealistic as N^2 handshakes must happen
 - EVS Engine
 - COReL



Relative Performance

http://www.cnds.jhu.edu/pub/papers/AT02_icdcs.pdf

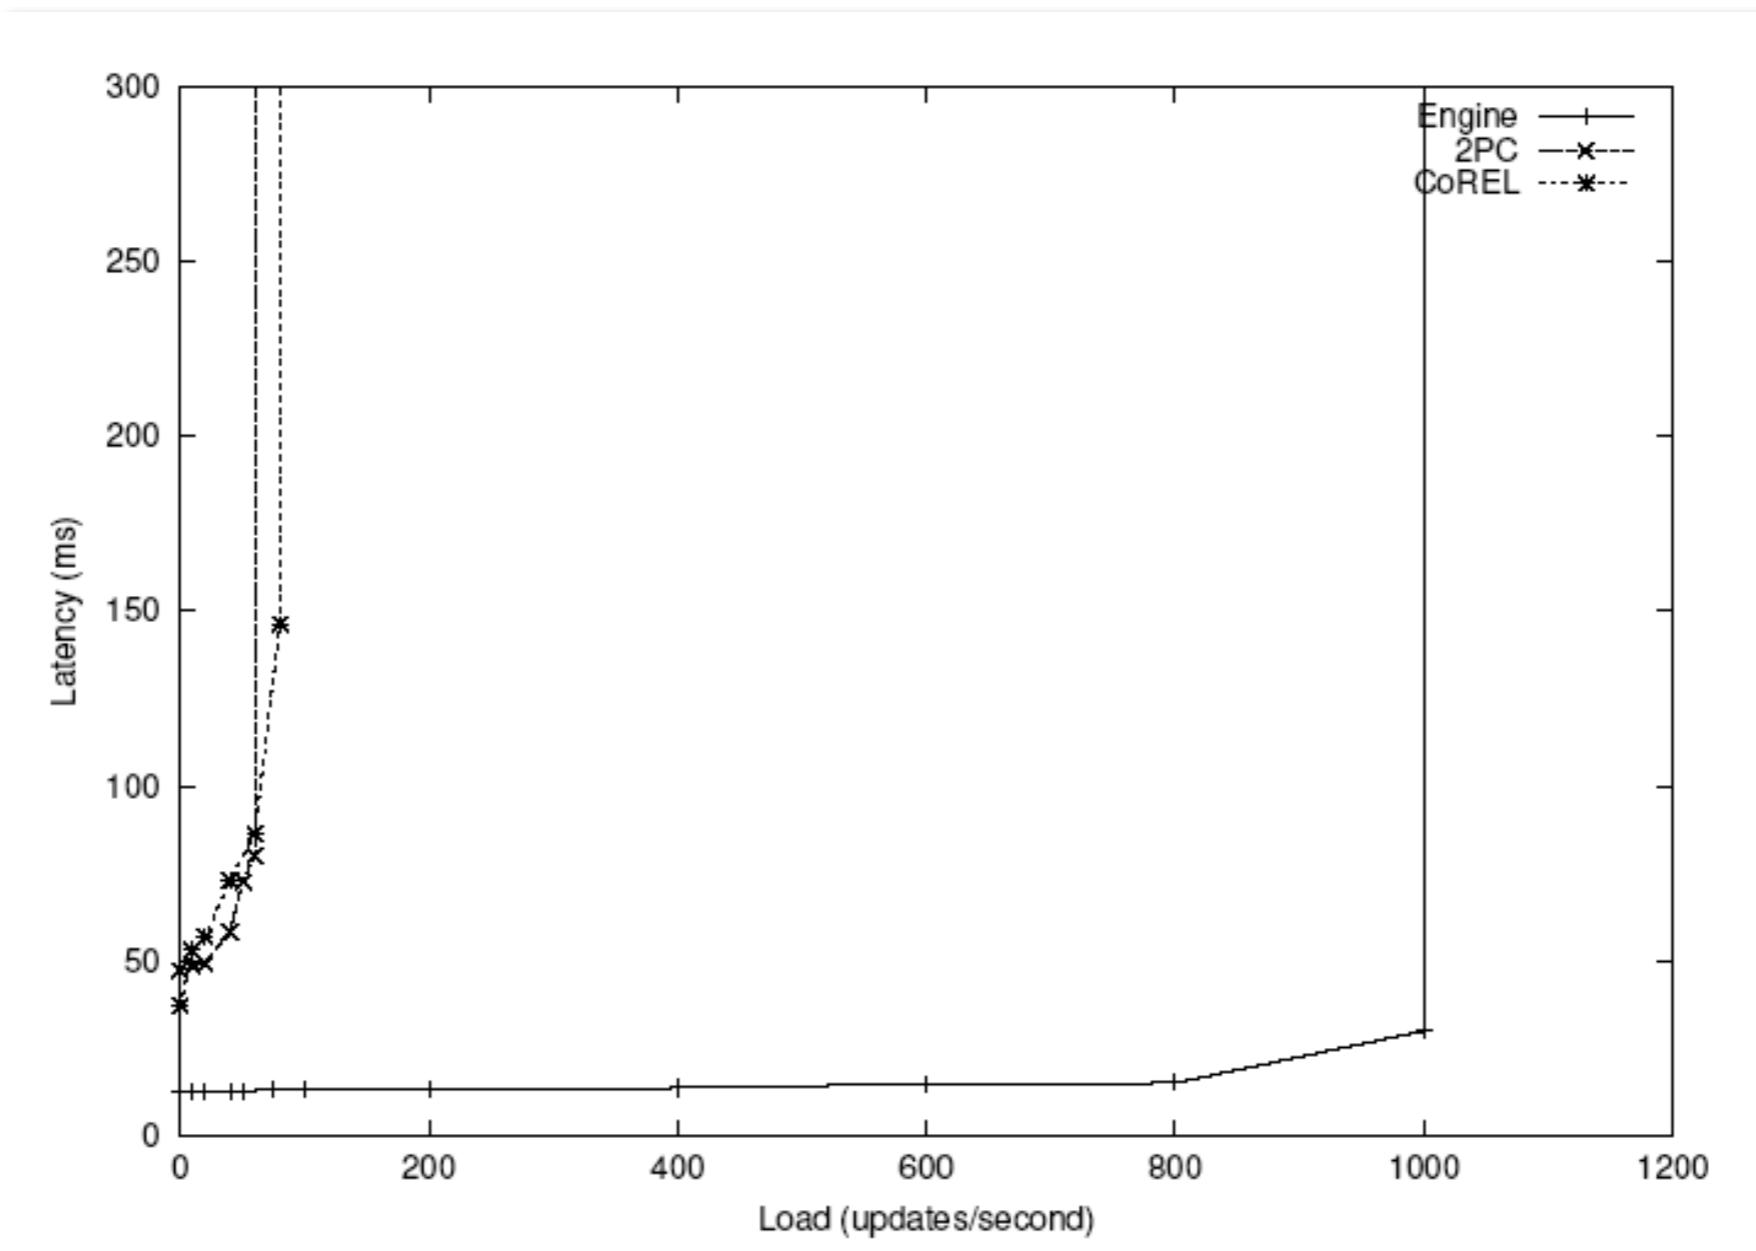
<http://www.cnds.jhu.edu/pub/papers/cnds-2002-4.pdf>



Relative Performance

http://www.cnds.jhu.edu/pub/papers/AT02_icdcs.pdf

<http://www.cnds.jhu.edu/pub/papers/cnnds-2002-4.pdf>



State of Affairs

- Multi-Master replication is a long way off
 - current implementors use 2PC
 - no enterprise offerings achieve EVS Engine performance
 - architectures that push databases hard aren't willing to cut performance for replication
 - multi-master is ready for architectures with low update rates that demand replication for data safety

- Master-Slave is ready for prime time
 - MySQL (native master-slave replication)
 - Oracle snapshots/materialized views

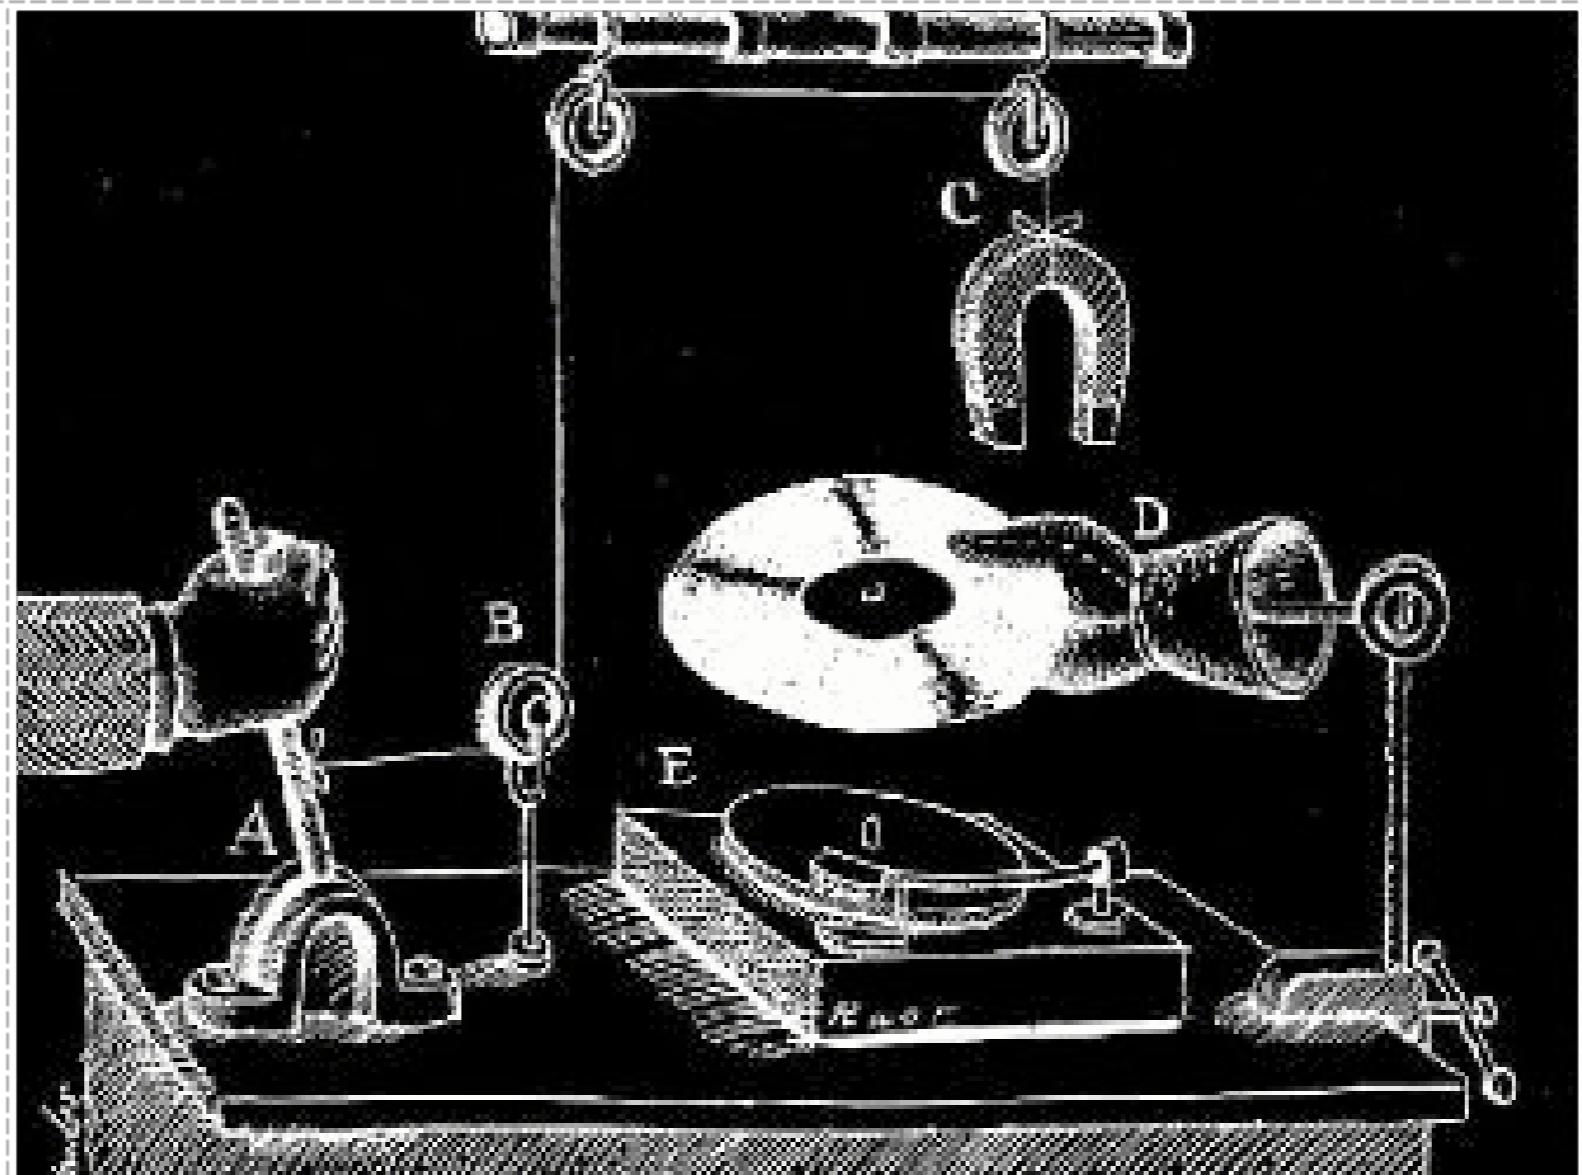


News Site Revisited

- ➊ Replicate the database on each web server
 - ➌ Oracle on each web server
 - ➌ replicate the needed tables
 - ➌ certainly doesn't scale financially

- ➋ If the site used MySQL...
 - ➌ zero capital investment
 - ➌ news items don't need to be cached in PHP pages
 - ➌ legitimizes more intense queries in live site pages





The Right Tool For The Job



Who's Online?

a real-world example

- A “service” requires who’s online info
- Users that have loaded an object within x minutes
- Need to know the last page the user hit
- The “service” is exposed throughout the site



Scalability Requirements

- $x = 30$ (minutes)
- 5000 hits/second
- 100,000 concurrent users
- Queries:
 - current users online (count)
 - current users on “this” page sorted by last access (limit 30)



Let's use a familiar tool

We use MySQL anyway

- We use MySQL to “drive” the site
- We are familiar with MySQL
- Queries are cake:
 - `select count(1) from recent_hits
where hitdate > SUBDATE(NOW(), INTERVAL 30 MINUTE)`
 - `select username, hitdate from recent_hits
where url = ? and hitdate > SUBDATE(NOW(), INTERVAL 30 MINUTE)
order by hitdate desc
limit 30`



Getting More Specific

- 100,000+ row table
 - assuming we sweep out stale data
- 5,000 replaces/second
 - indexes required on hitdate and url
- 1,000 queries/second
 - MySQL's query cache doesn't help at all, the updates invalidate it
 - Replaces cannot block queries
MyISAM is not an option, we use InnoDB
 - Both queries require a full table scan!



All in addition to the existing demands of the site!



Try Some Tests

- **On the development box
(Idle dual Xeon, SCSI drives)**
 - 1,400 replaces/second
 - 800 queries/second

- **On the production box
(dual Xeon, SCSI disk array w/ 1GB cache)**
 - 200 replaces/second
 - 150 queries/second

ARE YOU INSANE?!



Build a Custom Tool

- We need which urls/users/timestamp tracking
 - So... we need to add an “update” to each page
 - No... that won’t catch images, let’s use a mod_perl log hook.
 - Wait... we are already writing logs, let’s aggregate passively if we use mod_log_spread, we just need to add a subscriber
- Passive aggregation
 - can handle “bursty” traffic by lagging behind a bit
 - it can’t slow down the app!
- Pick a data structure
 - Multi-Indexed SkipList -- why?
 - Free “balancing” -- randomized
 - $O(\lg n)$ insertion, deletion, location
 - $O(1)$ popping (for culling expired sessions)
 - it precisely meets the requirements... and I like them.



Choose a Language

I like C... so sue me.

The concept:

- **Skiplist Index on:**
 - username
 - url,hitdate
 - hitdate
- **Single thread, event driven**
- **Receive messages from Spread:**
 - parse username, url, hitdate
 - delete from skiplist by username $O(\lg n)$
 - insert into skiplist $O(\lg n)$
 - pop the end of the skiplist of any hitdate > 30 minutes $O(1)$
- **Receive client requests**
 - Cardinality query is $O(1)$, single write()
 - Users on a url query is $O(\lg)$, $O(1)$ to fill out iovec, single writev()



The Result

- 6 hours worth of coding and testing
- 800 lines of C code (server)
 - use `libspread` and `libskiplist`
- 40 lines of perl (client module for web app)
- On commodity hardware (\$2k box)
 - ~80,000 inserts/second
 - more hits than we'll ever see!
 - ~100,000,000 counts/second -- not including `write()`
 - ~500,000 users for urls/second -- not including `writev()`

That'll do.





The Right Tool For The Job

- MySQL is a great tool
 - it is used to drive the rest of the site... spectacularly

- Using it for this project would:
 - wasted valuable resources in the architecture
 - saved a few hours of work
 - allowed you to not use your brain





This Image
Intentionally
Left Blank

Conclusion



Scalable Internet Architectures

- Building them just isn't that hard... if you
 - carefully analyze the problem at hand
 - don't make sloppy or rash technical decisions
 - always think like a computer scientist





Thank You

OmniTI Computer Consulting, Inc.

We're hiring:

Smart, fun and driven people

Buy The Book!!!

